

Assignment 2 Results

Qiao Wang

1. Consider the Cats & Dogs example. Start initially with a training sample of 1000, a validation sample of 500, and a test sample of 500 (like in the text). Use any technique to reduce overfitting and improve performance in developing a network that you train from scratch. What performance did you achieve?

Training a convnet from scratch with training sample of **1000**, a validation sample of 500, and a test sample of 500. Use augmentation and dropout(0.4), the test results are: **loss: 0.5437 - accuracy: 0.8490.**

2. Increase your training sample size. You may pick any amount. Keep the validation and test samples the same as above. Optimize your network (again training from scratch). What performance did you achieve?

Training a convnet from scratch with training sample of **3000**, a validation sample of 500, and a test sample of 500. Use augmentation and dropout(0.5), the test result is: **loss: 0.3999 - accuracy: 0.8868**

3. Now change your training sample so that you achieve better performance than those from Steps 1 and 2. This sample size may be larger, or smaller than those in the previous steps. The objective is to find the ideal training sample size to get best prediction results.

Training a convnet from scratch with training sample of **5000**, a validation sample of 500, and a test sample of 500. Use augmentation and dropout(0.5), the test results is: **loss: 0.2071 - accuracy: 0.9150.**

Training a convnet from scratch with training sample of **8000**, a validation sample of 500, and a test sample of 500. Use augmentation and dropout(0.5), the test results is: **loss: 0.2239 - accuracy: 0.9070.**

First, I increase the sample size to **5000**, the test result is: **loss: 0.2071 - accuracy: 0.9150.**

Then, I increase the sample size to **8000**, and we have a test result of **loss: 0.2239 - accuracy: 0.9070**, which is worse than when sample size is 5000.

Therefore, the best result I can get when training a convnet from scratch is: **loss: 0.2071 - accuracy: 0.9150, which is when the training sample size is 5000.**

4. Repeat Steps 1-3, but now using a pretrained network. The sample sizes you use in Steps 2 and 3 for the pretrained network may be the same or different from those using the network where you trained from scratch. Again, use any and all optimization techniques to get the best performance.

Using a pretrained network with training sample of **1000**, a validation sample of 500, and a test sample of 500. Use augmentation and dropout(0.5), the test results are: **loss: 1.5191 - accuracy: 0.9810**

Using a pretrained network with training sample of **3000**, a validation sample of 500, and a test sample of 500. Use augmentation and dropout(0.5), the test results are: **loss: 0.9459 - accuracy: 0.9790**

Using a pretrained network with training sample of **5000**, a validation sample of 500, and a test sample of 500. Use augmentation and dropout(0.5), the test results are: **loss: 0.3557 - accuracy: 0.9810**

Write a report summarizing your findings. What is the relationship between training sample size and choice of network?

Summary:

	training size	validation size	test size	augmentation	dropout	loss	accuracy
training from scratch	1000	500	500	yes	0.4	0.5437	0.849
	3000	500	500	yes	0.5	0.3999	0.8868
	5000	500	500	yes	0.5	0.2071	0.915
	8000	500	500	yes	0.5	0.2239	0.907
pretrained network	1000	500	500	yes	0.5	1.5191	0.981
	3000	500	500	yes	0.5	0.9459	0.979
	5000	500	500	yes	0.5	0.3557	0.981

As we see from the results above, using convnets and train from scratch could get a relatively good result for a small dataset. However, there are some overfitting problems. We can solve this problem to some degree by using techniques like dropout, augmentation, etc. When we set sample size to 1000, the accuracy is very low. As the sample size increases, the loss decreases and accuracy increases. When the sample size increases to 5000, the accuracy reaches peak and then decrease as sample size reach 8000.

When we use a pretrained network, the results of accuracy are all around 0.9800, although the loss reduces when the sample size increases. Therefore, in this case, the accuracy of the model doesn't change much, and the loss decreases as the sample size increases.

Training a convnet from scratch with training sample of 1000, a validation sample of 500, and a test sample of 500

Downloading the data

```
!unzip -qq '/fs/ess/PGS0333/BA_64061_KSU/data/dogs-vs-cats.zip'
```

 In [1]:

```
!unzip -qq train.zip
```

 In [2]:

Copying images to training, validation, and test directories

```
import os, shutil, pathlib
```

 In [3]:

```
original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index,
end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=1000)
make_subset("validation", start_index=1000, end_index=1500)
make_subset("test", start_index=1500, end_index=2000)
```

Building the model

Instantiating a small convnet for dogs vs. cats classification

```
from tensorflow import keras
from tensorflow.keras import layers

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
```

 In [4]:

```
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

In [5]:

```
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 180, 180, 3)]	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 256)	0
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590080
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1)	12545
=====		
Total params: 991,041		
Trainable params: 991,041		
Non-trainable params: 0		

Configuring the model for training

In [6]:

```
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

Data preprocessing

Using image_dataset_from_directory to read images

In [7]:

```
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)

Found 2000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
```

In [8]:

```
import numpy as np
import tensorflow as tf
random_numbers = np.random.normal(size=(1000, 16))
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)
```

In [9]:

```
for i, element in enumerate(dataset):
    print(element.shape)
    if i >= 2:
        break

(16,)
(16,)
(16,)
```

In [10]:

```
batched_dataset = dataset.batch(32)
for i, element in enumerate(batched_dataset):
    print(element.shape)
    if i >= 2:
        break

(32, 16)
(32, 16)
(32, 16)
```

In [11]:

```
reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))
for i, element in enumerate(reshaped_dataset):
    print(element.shape)
    if i >= 2:
        break

(4, 4)
(4, 4)
(4, 4)
```

Displaying the shapes of the data and labels yielded by the Dataset

In [12]:

```
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    break
```

data batch shape: (32, 180, 180, 3)

labels batch shape: (32,)

Fitting the model using a Dataset

In [13]:

```
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)

Epoch 1/30
63/63 [=====] - 5s 38ms/step - loss: 0.7281 - accuracy: 0.5025 - val_loss: 0.7877 - val_accuracy: 0.5000
Epoch 2/30
63/63 [=====] - 2s 30ms/step - loss: 0.7313 - accuracy: 0.5295 - val_loss: 0.6778 - val_accuracy: 0.6290
Epoch 3/30
63/63 [=====] - 2s 28ms/step - loss: 0.6879 - accuracy: 0.6005 - val_loss: 0.7451 - val_accuracy: 0.5220
Epoch 4/30
63/63 [=====] - 2s 28ms/step - loss: 0.6479 - accuracy: 0.6325 - val_loss: 0.7421 - val_accuracy: 0.5580
Epoch 5/30
63/63 [=====] - 2s 28ms/step - loss: 0.6098 - accuracy: 0.6705 - val_loss: 0.6288 - val_accuracy: 0.6430
Epoch 6/30
63/63 [=====] - 2s 28ms/step - loss: 0.5627 - accuracy: 0.7020 - val_loss: 0.6872 - val_accuracy: 0.6230
Epoch 7/30
63/63 [=====] - 2s 28ms/step - loss: 0.5525 - accuracy: 0.7360 - val_loss: 0.9677 - val_accuracy: 0.6000
Epoch 8/30
63/63 [=====] - 2s 27ms/step - loss: 0.4833 - accuracy: 0.7745 - val_loss: 0.8117 - val_accuracy: 0.6070
Epoch 9/30
63/63 [=====] - 2s 28ms/step - loss: 0.4395 - accuracy: 0.7955 - val_loss: 0.6178 - val_accuracy: 0.6920
Epoch 10/30
63/63 [=====] - 2s 28ms/step - loss: 0.3887 - accuracy: 0.8390 - val_loss: 0.5901 - val_accuracy: 0.7150
Epoch 11/30
```

63/63 [=====] - 2s 28ms/step - loss: 0.3209 - accuracy: 0.8645 - val_loss: 0.6284 - val_accuracy: 0.7260
Epoch 12/30
63/63 [=====] - 2s 28ms/step - loss: 0.2802 - accuracy: 0.8885 - val_loss: 0.7380 - val_accuracy: 0.7420
Epoch 13/30
63/63 [=====] - 2s 28ms/step - loss: 0.2119 - accuracy: 0.9175 - val_loss: 0.6878 - val_accuracy: 0.7290
Epoch 14/30
63/63 [=====] - 2s 28ms/step - loss: 0.1710 - accuracy: 0.9315 - val_loss: 0.8529 - val_accuracy: 0.7380
Epoch 15/30
63/63 [=====] - 2s 29ms/step - loss: 0.1394 - accuracy: 0.9445 - val_loss: 0.9875 - val_accuracy: 0.7410
Epoch 16/30
63/63 [=====] - 2s 28ms/step - loss: 0.1168 - accuracy: 0.9540 - val_loss: 0.8968 - val_accuracy: 0.7720
Epoch 17/30
63/63 [=====] - 2s 28ms/step - loss: 0.0858 - accuracy: 0.9690 - val_loss: 1.1203 - val_accuracy: 0.7760
Epoch 18/30
63/63 [=====] - 2s 28ms/step - loss: 0.1041 - accuracy: 0.9700 - val_loss: 1.1405 - val_accuracy: 0.7580
Epoch 19/30
63/63 [=====] - 2s 29ms/step - loss: 0.0519 - accuracy: 0.9825 - val_loss: 1.6991 - val_accuracy: 0.7490
Epoch 20/30
63/63 [=====] - 2s 28ms/step - loss: 0.0868 - accuracy: 0.9675 - val_loss: 1.4378 - val_accuracy: 0.7130
Epoch 21/30
63/63 [=====] - 2s 29ms/step - loss: 0.0478 - accuracy: 0.9830 - val_loss: 1.4454 - val_accuracy: 0.7510
Epoch 22/30
63/63 [=====] - 2s 28ms/step - loss: 0.0584 - accuracy: 0.9800 - val_loss: 1.3630 - val_accuracy: 0.7520
Epoch 23/30
63/63 [=====] - 2s 28ms/step - loss: 0.0501 - accuracy: 0.9835 - val_loss: 1.8856 - val_accuracy: 0.7040
Epoch 24/30
63/63 [=====] - 2s 29ms/step - loss: 0.0361 - accuracy: 0.9890 - val_loss: 1.8424 - val_accuracy: 0.7590
Epoch 25/30
63/63 [=====] - 2s 28ms/step - loss: 0.0699 - accuracy: 0.9800 - val_loss: 1.5784 - val_accuracy: 0.7460
Epoch 26/30
63/63 [=====] - 2s 28ms/step - loss: 0.0507 - accuracy: 0.9835 - val_loss: 1.7306 - val_accuracy: 0.7450
Epoch 27/30
63/63 [=====] - 2s 29ms/step - loss: 0.0408 - accuracy: 0.9885 - val_loss: 1.7925 - val_accuracy: 0.7320
Epoch 28/30
63/63 [=====] - 2s 28ms/step - loss: 0.0331 - accuracy: 0.9890 - val_loss: 1.9974 - val_accuracy: 0.7360

Epoch 29/30

63/63 [=====] - 2s 27ms/step - loss: 0.0544 - accuracy: 0.9860 - val_loss: 2.1735 - val_accuracy: 0.7330


Epoch 30/30


63/63 [=====] - 2s 28ms/step - loss: 0.0622 - accuracy: 0.9845 - val_loss: 2.3797 - val_accuracy: 0.7230

Displaying curves of loss and accuracy during training

In [14]:

```
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```





Evaluating the model on the test set

In [15]:

```
test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 14ms/step - loss: 0.6620 - accuracy: 0.6860
Test accuracy: 0.686
```

Using data augmentation

Define a data augmentation stage to add to an image model

In [16]:

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
```

Displaying some randomly augmented training images

In [17]:

```
plt.figure(figsize=(10, 10))
for images, _ in train_dataset.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_images[0].numpy().astype("uint8"))
```

```
plt.axis("off")
```



Defining a new convnet that includes image augmentation, and dropout

```
inputs = keras.Input(shape=(180, 180, 3))  
x = data_augmentation(inputs)  
x = layers.Rescaling(1./255)(x)
```

In [18]:

```

x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.4)(x)      # change the dropout from 0.5 to 0.4
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

Training the regularized convnet

In [19]:

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=80, # change epochs from 100 to 80
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/80
63/63 [=====] - 3s 31ms/step - loss: 0.7067 - accuracy: 0.5165 - val_loss: 0.6852 - val_accuracy: 0.5560
Epoch 2/80
63/63 [=====] - 2s 29ms/step - loss: 0.7041 - accuracy: 0.5600 - val_loss: 0.6838 - val_accuracy: 0.5500
Epoch 3/80
63/63 [=====] - 2s 29ms/step - loss: 0.6729 - accuracy: 0.6110 - val_loss: 0.6920 - val_accuracy: 0.5460
Epoch 4/80
63/63 [=====] - 2s 29ms/step - loss: 0.6557 - accuracy: 0.6225 - val_loss: 0.6604 - val_accuracy: 0.5790
Epoch 5/80
63/63 [=====] - 2s 29ms/step - loss: 0.6495 - accuracy: 0.6425 - val_loss: 0.6610 - val_accuracy: 0.6140
Epoch 6/80
63/63 [=====] - 2s 29ms/step - loss: 0.6109 - accuracy: 0.6540 - val_loss: 1.1658 - val_accuracy: 0.5730
Epoch 7/80
63/63 [=====] - 2s 29ms/step - loss: 0.6154 - accuracy: 0.6795 - val_loss: 0.5736 - val_accuracy: 0.6950

Epoch 8/80
63/63 [=====] - 2s 29ms/step - loss: 0.5897 - accuracy: 0.7035 - val_loss: 0.6279 - val_accuracy: 0.6810
Epoch 9/80
63/63 [=====] - 2s 30ms/step - loss: 0.5802 - accuracy: 0.7040 - val_loss: 0.6419 - val_accuracy: 0.6910
Epoch 10/80
63/63 [=====] - 2s 30ms/step - loss: 0.5710 - accuracy: 0.7035 - val_loss: 0.5431 - val_accuracy: 0.7350
Epoch 11/80
63/63 [=====] - 2s 29ms/step - loss: 0.5470 - accuracy: 0.7240 - val_loss: 0.6522 - val_accuracy: 0.6880
Epoch 12/80
63/63 [=====] - 2s 29ms/step - loss: 0.5214 - accuracy: 0.7425 - val_loss: 0.5042 - val_accuracy: 0.7430
Epoch 13/80
63/63 [=====] - 2s 28ms/step - loss: 0.5125 - accuracy: 0.7565 - val_loss: 0.6554 - val_accuracy: 0.7340
Epoch 14/80
63/63 [=====] - 2s 28ms/step - loss: 0.5010 - accuracy: 0.7500 - val_loss: 0.5015 - val_accuracy: 0.7540
Epoch 15/80
63/63 [=====] - 2s 28ms/step - loss: 0.4955 - accuracy: 0.7630 - val_loss: 0.6987 - val_accuracy: 0.6910
Epoch 16/80
63/63 [=====] - 2s 29ms/step - loss: 0.4789 - accuracy: 0.7775 - val_loss: 0.5838 - val_accuracy: 0.7250
Epoch 17/80
63/63 [=====] - 2s 29ms/step - loss: 0.4715 - accuracy: 0.7775 - val_loss: 0.4991 - val_accuracy: 0.7650
Epoch 18/80
63/63 [=====] - 2s 29ms/step - loss: 0.4562 - accuracy: 0.7945 - val_loss: 0.5710 - val_accuracy: 0.7350
Epoch 19/80
63/63 [=====] - 2s 28ms/step - loss: 0.4658 - accuracy: 0.7915 - val_loss: 0.4418 - val_accuracy: 0.7900
Epoch 20/80
63/63 [=====] - 2s 29ms/step - loss: 0.4599 - accuracy: 0.7870 - val_loss: 0.4657 - val_accuracy: 0.7790
Epoch 21/80
63/63 [=====] - 2s 29ms/step - loss: 0.4483 - accuracy: 0.7930 - val_loss: 0.5815 - val_accuracy: 0.7030
Epoch 22/80
63/63 [=====] - 2s 28ms/step - loss: 0.4382 - accuracy: 0.7940 - val_loss: 0.6202 - val_accuracy: 0.7490
Epoch 23/80
63/63 [=====] - 2s 29ms/step - loss: 0.4339 - accuracy: 0.8005 - val_loss: 0.4995 - val_accuracy: 0.7800
Epoch 24/80
63/63 [=====] - 2s 29ms/step - loss: 0.4112 - accuracy: 0.8145 - val_loss: 0.4665 - val_accuracy: 0.7950
Epoch 25/80

63/63 [=====] - 2s 29ms/step - loss: 0.3912 - accuracy: 0.8330 - val_loss: 0.5159 - val_accuracy: 0.7550
Epoch 26/80
63/63 [=====] - 2s 28ms/step - loss: 0.3978 - accuracy: 0.8225 - val_loss: 0.4460 - val_accuracy: 0.8110
Epoch 27/80
63/63 [=====] - 2s 29ms/step - loss: 0.3964 - accuracy: 0.8155 - val_loss: 0.4716 - val_accuracy: 0.7960
Epoch 28/80
63/63 [=====] - 2s 29ms/step - loss: 0.3859 - accuracy: 0.8260 - val_loss: 0.6183 - val_accuracy: 0.7300
Epoch 29/80
63/63 [=====] - 2s 29ms/step - loss: 0.3803 - accuracy: 0.8360 - val_loss: 0.5046 - val_accuracy: 0.7730
Epoch 30/80
63/63 [=====] - 2s 29ms/step - loss: 0.3758 - accuracy: 0.8315 - val_loss: 0.7977 - val_accuracy: 0.7130
Epoch 31/80
63/63 [=====] - 2s 28ms/step - loss: 0.3527 - accuracy: 0.8375 - val_loss: 0.6236 - val_accuracy: 0.8000
Epoch 32/80
63/63 [=====] - 2s 29ms/step - loss: 0.3576 - accuracy: 0.8425 - val_loss: 0.5660 - val_accuracy: 0.7640
Epoch 33/80
63/63 [=====] - 2s 29ms/step - loss: 0.3353 - accuracy: 0.8540 - val_loss: 0.5750 - val_accuracy: 0.8180
Epoch 34/80
63/63 [=====] - 2s 28ms/step - loss: 0.3402 - accuracy: 0.8575 - val_loss: 0.5400 - val_accuracy: 0.7850
Epoch 35/80
63/63 [=====] - 2s 29ms/step - loss: 0.3140 - accuracy: 0.8650 - val_loss: 0.5461 - val_accuracy: 0.7930
Epoch 36/80
63/63 [=====] - 2s 29ms/step - loss: 0.3298 - accuracy: 0.8655 - val_loss: 0.4594 - val_accuracy: 0.7960
Epoch 37/80
63/63 [=====] - 2s 29ms/step - loss: 0.3310 - accuracy: 0.8670 - val_loss: 0.6904 - val_accuracy: 0.7680
Epoch 38/80
63/63 [=====] - 2s 28ms/step - loss: 0.3049 - accuracy: 0.8745 - val_loss: 0.5247 - val_accuracy: 0.8190
Epoch 39/80
63/63 [=====] - 2s 29ms/step - loss: 0.3203 - accuracy: 0.8685 - val_loss: 0.4621 - val_accuracy: 0.8020
Epoch 40/80
63/63 [=====] - 2s 29ms/step - loss: 0.2851 - accuracy: 0.8790 - val_loss: 0.7653 - val_accuracy: 0.7940
Epoch 41/80
63/63 [=====] - 2s 29ms/step - loss: 0.3136 - accuracy: 0.8660 - val_loss: 0.4930 - val_accuracy: 0.8060
Epoch 42/80
63/63 [=====] - 2s 28ms/step - loss: 0.3056 - accuracy: 0.8685 - val_loss: 0.4781 - val_accuracy: 0.8120

Epoch 43/80
63/63 [=====] - 2s 30ms/step - loss: 0.2810 - accuracy: 0.8810 - val_loss: 0.4700 - val_accuracy: 0.8290
Epoch 44/80
63/63 [=====] - 2s 30ms/step - loss: 0.2661 - accuracy: 0.8930 - val_loss: 1.2875 - val_accuracy: 0.6740
Epoch 45/80
63/63 [=====] - 2s 30ms/step - loss: 0.2815 - accuracy: 0.8910 - val_loss: 0.5469 - val_accuracy: 0.7930
Epoch 46/80
63/63 [=====] - 2s 29ms/step - loss: 0.2696 - accuracy: 0.8940 - val_loss: 0.8227 - val_accuracy: 0.7920
Epoch 47/80
63/63 [=====] - 2s 29ms/step - loss: 0.2462 - accuracy: 0.8980 - val_loss: 0.5565 - val_accuracy: 0.8330
Epoch 48/80
63/63 [=====] - 2s 29ms/step - loss: 0.2627 - accuracy: 0.8925 - val_loss: 0.5671 - val_accuracy: 0.7800
Epoch 49/80
63/63 [=====] - 2s 30ms/step - loss: 0.2446 - accuracy: 0.8915 - val_loss: 0.5453 - val_accuracy: 0.8010
Epoch 50/80
63/63 [=====] - 2s 29ms/step - loss: 0.2383 - accuracy: 0.8990 - val_loss: 0.4764 - val_accuracy: 0.8470
Epoch 51/80
63/63 [=====] - 2s 29ms/step - loss: 0.2692 - accuracy: 0.8920 - val_loss: 0.4775 - val_accuracy: 0.8160
Epoch 52/80
63/63 [=====] - 2s 28ms/step - loss: 0.2244 - accuracy: 0.9120 - val_loss: 0.7375 - val_accuracy: 0.8030
Epoch 53/80
63/63 [=====] - 2s 28ms/step - loss: 0.2228 - accuracy: 0.9065 - val_loss: 0.5973 - val_accuracy: 0.8110
Epoch 54/80
63/63 [=====] - 2s 29ms/step - loss: 0.2456 - accuracy: 0.9010 - val_loss: 0.5144 - val_accuracy: 0.8180
Epoch 55/80
63/63 [=====] - 2s 28ms/step - loss: 0.2151 - accuracy: 0.9180 - val_loss: 0.7244 - val_accuracy: 0.8210
Epoch 56/80
63/63 [=====] - 2s 29ms/step - loss: 0.2265 - accuracy: 0.9105 - val_loss: 0.5406 - val_accuracy: 0.8110
Epoch 57/80
63/63 [=====] - 2s 29ms/step - loss: 0.2274 - accuracy: 0.9110 - val_loss: 0.5016 - val_accuracy: 0.8160
Epoch 58/80
63/63 [=====] - 2s 29ms/step - loss: 0.2249 - accuracy: 0.9165 - val_loss: 0.4593 - val_accuracy: 0.8380
Epoch 59/80
63/63 [=====] - 2s 29ms/step - loss: 0.2030 - accuracy: 0.9265 - val_loss: 0.6098 - val_accuracy: 0.8050
Epoch 60/80

63/63 [=====] - 3s 39ms/step - loss: 0.2300 - accuracy: 0.9165 - val_loss: 0.7569 - val_accuracy: 0.8280
Epoch 61/80
63/63 [=====] - 2s 31ms/step - loss: 0.2432 - accuracy: 0.9115 - val_loss: 0.4634 - val_accuracy: 0.8470
Epoch 62/80
63/63 [=====] - 2s 31ms/step - loss: 0.1986 - accuracy: 0.9270 - val_loss: 0.5855 - val_accuracy: 0.8290
Epoch 63/80
63/63 [=====] - 2s 31ms/step - loss: 0.1942 - accuracy: 0.9320 - val_loss: 0.5501 - val_accuracy: 0.8410
Epoch 64/80
63/63 [=====] - 2s 30ms/step - loss: 0.2051 - accuracy: 0.9255 - val_loss: 0.5786 - val_accuracy: 0.8270
Epoch 65/80
63/63 [=====] - 2s 30ms/step - loss: 0.1830 - accuracy: 0.9305 - val_loss: 0.6350 - val_accuracy: 0.8230
Epoch 66/80
63/63 [=====] - 2s 30ms/step - loss: 0.2005 - accuracy: 0.9205 - val_loss: 0.5401 - val_accuracy: 0.8370
Epoch 67/80
63/63 [=====] - 2s 31ms/step - loss: 0.1982 - accuracy: 0.9245 - val_loss: 0.6137 - val_accuracy: 0.8020
Epoch 68/80
63/63 [=====] - 2s 31ms/step - loss: 0.2312 - accuracy: 0.9130 - val_loss: 0.5874 - val_accuracy: 0.8300
Epoch 69/80
63/63 [=====] - 2s 31ms/step - loss: 0.1951 - accuracy: 0.9260 - val_loss: 0.4314 - val_accuracy: 0.8550
Epoch 70/80
63/63 [=====] - 2s 31ms/step - loss: 0.1619 - accuracy: 0.9320 - val_loss: 0.7138 - val_accuracy: 0.8430
Epoch 71/80
63/63 [=====] - 2s 31ms/step - loss: 0.2078 - accuracy: 0.9200 - val_loss: 0.4732 - val_accuracy: 0.8670
Epoch 72/80
63/63 [=====] - 2s 31ms/step - loss: 0.1862 - accuracy: 0.9265 - val_loss: 1.2674 - val_accuracy: 0.7960
Epoch 73/80
63/63 [=====] - 2s 30ms/step - loss: 0.1802 - accuracy: 0.9345 - val_loss: 0.5652 - val_accuracy: 0.8480
Epoch 74/80
63/63 [=====] - 2s 31ms/step - loss: 0.1870 - accuracy: 0.9320 - val_loss: 0.6469 - val_accuracy: 0.8540
Epoch 75/80
63/63 [=====] - 2s 30ms/step - loss: 0.1809 - accuracy: 0.9335 - val_loss: 0.7291 - val_accuracy: 0.8450
Epoch 76/80
63/63 [=====] - 2s 32ms/step - loss: 0.1844 - accuracy: 0.9235 - val_loss: 0.5942 - val_accuracy: 0.8370
Epoch 77/80
63/63 [=====] - 2s 31ms/step - loss: 0.1893 - accuracy: 0.9335 - val_loss: 0.5588 - val_accuracy: 0.8310

```
Epoch 78/80
63/63 [=====] - 2s 30ms/step - loss: 0.1505 - accuracy: 0.9455 - val_loss: 0.5778 - val_accuracy: 0.8330
Epoch 79/80
63/63 [=====] - 2s 30ms/step - loss: 0.1783 - accuracy: 0.9325 - val_loss: 0.8527 - val_accuracy: 0.8380
Epoch 80/80
63/63 [=====] - 2s 30ms/step - loss: 0.1884 - accuracy: 0.9310 - val_loss: 0.7090 - val_accuracy: 0.8270
```

Evaluating the model on the test set

In [20]:

```
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 15ms/step - loss: 0.5437 - accuracy: 0.8490
Test accuracy: 0.849
```


Training a convnet from scratch with training sample of 3000, a validation sample of 500, and a test sample of 500

Downloading the data

```
#!/unzip -qq '/fs/ess/PGS0333/BA_64061_KSU/data/dogs-vs-cats.zip'
```

In [1]:

```
#!/unzip -qq train.zip
```

In [2]:

Copying images to training, validation, and test directories

```
import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index,
end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=2999)
make_subset("validation", start_index=3000, end_index=3499)
make_subset("test", start_index=3500, end_index=3999)
```

In [3]:

Building the model

Instantiating a small convnet for dogs vs. cats classification

```
from tensorflow import keras
from tensorflow.keras import layers

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
```

In [4]:

```
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

In [5]:

```
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 180, 180, 3)]	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 256)	0
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590080
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1)	12545
=====		
Total params: 991,041		
Trainable params: 991,041		
Non-trainable params: 0		

Configuring the model for training

In [6]:

```
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

Data preprocessing

Using image_dataset_from_directory to read images

In [7]:

```
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)

Found 5998 files belonging to 2 classes.
Found 998 files belonging to 2 classes.
Found 998 files belonging to 2 classes.
```

In [8]:

```
import numpy as np
import tensorflow as tf
random_numbers = np.random.normal(size=(1000, 16))
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)
```

In [9]:

```
for i, element in enumerate(dataset):
    print(element.shape)
    if i >= 2:
        break

(16,)
(16,)
(16,)
```

In [10]:

```
batched_dataset = dataset.batch(32)
for i, element in enumerate(batched_dataset):
    print(element.shape)
    if i >= 2:
        break

(32, 16)
(32, 16)
(32, 16)
```

In [11]:

```
reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))
for i, element in enumerate(reshaped_dataset):
    print(element.shape)
    if i >= 2:
        break

(4, 4)
(4, 4)
(4, 4)
```

Displaying the shapes of the data and labels yielded by the Dataset

In [12]:

```
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    break
```

data batch shape: (32, 180, 180, 3)

labels batch shape: (32,)

Fitting the model using a Dataset

In [13]:

```
"""
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)
"""
```

Out[13]:

```
'\ncallbacks = [\n    keras.callbacks.ModelCheckpoint(\n        filepath="convnet_from_scratch.keras",\n        save_best_only=True,\n        monitor="val_loss")\n]\n\nhistory = model.fit(\n    train_dataset,\n    epochs=30,\n    validation_data=validation_dataset,\n    callbacks=callbacks)\n'
```

Displaying curves of loss and accuracy during training

In [14]:

```
"""
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
"""
```

Out[14]:

```
'\nimport matplotlib.pyplot as plt\n\naccuracy = history.history["accuracy"]\nval_accuracy = history.history["val_accuracy"]\n\nloss = history.history["loss"]
```

```
\nval_loss = history.history["val_loss"]\nepochs = range(1, len(accuracy) + 1)\n\nplt.plot(epochs, accuracy, "bo", label="Training accuracy")\nplt.plot(epochs, val_accuracy, "b", label="Validation accuracy")\nplt.title("Training and validation accuracy")\nplt.legend()\nplt.figure()\nplt.plot(epochs, loss, "bo", label="Training loss")\nplt.plot(epochs, val_loss, "b", label="Validation loss")\nplt.title("Training and validation loss")\nplt.legend()\nplt.show()\n'
```

Evaluating the model on the test set

In [15]:

```
"""\ntest_model = keras.models.load_model("convnet_from_scratch.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n"""
```

Out[15]:

```
'\ntest_model = keras.models.load_model("convnet_from_scratch.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n'
```

Using data augmentation

Define a data augmentation stage to add to an image model

In [16]:

```
data_augmentation = keras.Sequential(\n    [\n        layers.RandomFlip("horizontal"),\n        layers.RandomRotation(0.1),\n        layers.RandomZoom(0.2),\n    ]\n)
```

Displaying some randomly augmented training images

In [17]:

```
"""\nplt.figure(figsize=(10, 10))\nfor images, _ in train_dataset.take(1):\n    for i in range(9):\n        augmented_images = data_augmentation(images)\n        ax = plt.subplot(3, 3, i + 1)\n        plt.imshow(augmented_images[0].numpy().astype("uint8"))\n        plt.axis("off")\n"""\n
```

Out[17]:

```
'\nplt.figure(figsize=(10, 10))\nfor images, _ in train_dataset.take(1):\n    for i in range(9):\n        augmented_images = data_augmentation(images)\n        ax = plt.subplot(3, 3, i + 1)\n        plt.imshow(augmented_images[0].numpy().astype("uint8"))\n        plt.axis("off")\n'
```

Defining a new convnet that includes image augmentation, regularization and dropout

In [18]:

```
inputs = keras.Input(shape=(180, 180, 3))
```

```

x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

Training the regularized convnet

In [19]:

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=80,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/80
188/188 [=====] - 9s 28ms/step - loss: 0.7013 - accuracy: 0.5367 - val_loss: 0.6745 - val_accuracy: 0.6222
Epoch 2/80
188/188 [=====] - 5s 25ms/step - loss: 0.6622 - accuracy: 0.6307 - val_loss: 0.6223 - val_accuracy: 0.6774
Epoch 3/80
188/188 [=====] - 5s 25ms/step - loss: 0.6244 - accuracy: 0.6557 - val_loss: 0.5671 - val_accuracy: 0.7194
Epoch 4/80
188/188 [=====] - 5s 25ms/step - loss: 0.5998 - accuracy: 0.6812 - val_loss: 0.6425 - val_accuracy: 0.6683
Epoch 5/80
188/188 [=====] - 5s 24ms/step - loss: 0.5709 - accuracy: 0.7062 - val_loss: 0.5270 - val_accuracy: 0.7545
Epoch 6/80
188/188 [=====] - 5s 24ms/step - loss: 0.5397 - accuracy: 0.7249 - val_loss: 0.4980 - val_accuracy: 0.7625
Epoch 7/80

188/188 [=====] - 5s 24ms/step - loss: 0.5218 - accuracy: 0.7487 - val_loss: 0.4215 - val_accuracy: 0.8146
Epoch 8/80
188/188 [=====] - 5s 24ms/step - loss: 0.4977 - accuracy: 0.7598 - val_loss: 0.4423 - val_accuracy: 0.7976
Epoch 9/80
188/188 [=====] - 5s 25ms/step - loss: 0.4810 - accuracy: 0.7749 - val_loss: 0.4341 - val_accuracy: 0.8096
Epoch 10/80
188/188 [=====] - 5s 24ms/step - loss: 0.4671 - accuracy: 0.7798 - val_loss: 0.4010 - val_accuracy: 0.8196
Epoch 11/80
188/188 [=====] - 5s 24ms/step - loss: 0.4539 - accuracy: 0.7893 - val_loss: 0.3904 - val_accuracy: 0.8096
Epoch 12/80
188/188 [=====] - 5s 24ms/step - loss: 0.4378 - accuracy: 0.8018 - val_loss: 0.3867 - val_accuracy: 0.8196
Epoch 13/80
188/188 [=====] - 5s 24ms/step - loss: 0.4131 - accuracy: 0.8129 - val_loss: 0.3734 - val_accuracy: 0.8297
Epoch 14/80
188/188 [=====] - 5s 24ms/step - loss: 0.4097 - accuracy: 0.8184 - val_loss: 0.3403 - val_accuracy: 0.8637
Epoch 15/80
188/188 [=====] - 5s 25ms/step - loss: 0.4037 - accuracy: 0.8268 - val_loss: 0.3120 - val_accuracy: 0.8637
Epoch 16/80
188/188 [=====] - 5s 24ms/step - loss: 0.3826 - accuracy: 0.8336 - val_loss: 0.4983 - val_accuracy: 0.7936
Epoch 17/80
188/188 [=====] - 5s 25ms/step - loss: 0.3803 - accuracy: 0.8348 - val_loss: 0.4090 - val_accuracy: 0.8537
Epoch 18/80
188/188 [=====] - 5s 25ms/step - loss: 0.3632 - accuracy: 0.8408 - val_loss: 0.3277 - val_accuracy: 0.8848
Epoch 19/80
188/188 [=====] - 5s 25ms/step - loss: 0.3537 - accuracy: 0.8503 - val_loss: 0.4235 - val_accuracy: 0.8317
Epoch 20/80
188/188 [=====] - 5s 24ms/step - loss: 0.3514 - accuracy: 0.8546 - val_loss: 0.2837 - val_accuracy: 0.8818
Epoch 21/80
188/188 [=====] - 5s 24ms/step - loss: 0.3369 - accuracy: 0.8550 - val_loss: 0.3131 - val_accuracy: 0.8828
Epoch 22/80
188/188 [=====] - 5s 25ms/step - loss: 0.3493 - accuracy: 0.8590 - val_loss: 0.3317 - val_accuracy: 0.8557
Epoch 23/80
188/188 [=====] - 5s 26ms/step - loss: 0.3231 - accuracy: 0.8615 - val_loss: 0.3820 - val_accuracy: 0.8457
Epoch 24/80
188/188 [=====] - 5s 24ms/step - loss: 0.3298 - accuracy: 0.8573 - val_loss: 0.4611 - val_accuracy: 0.7836

Epoch 25/80
188/188 [=====] - 5s 24ms/step - loss: 0.3255 - accuracy: 0.8560 - val_loss: 0.3417 - val_accuracy: 0.8717
Epoch 26/80
188/188 [=====] - 5s 25ms/step - loss: 0.3262 - accuracy: 0.8651 - val_loss: 0.2503 - val_accuracy: 0.9108
Epoch 27/80
188/188 [=====] - 5s 24ms/step - loss: 0.3250 - accuracy: 0.8646 - val_loss: 0.2798 - val_accuracy: 0.9068
Epoch 28/80
188/188 [=====] - 5s 24ms/step - loss: 0.3126 - accuracy: 0.8733 - val_loss: 0.3802 - val_accuracy: 0.8577
Epoch 29/80
188/188 [=====] - 5s 24ms/step - loss: 0.2980 - accuracy: 0.8746 - val_loss: 0.2611 - val_accuracy: 0.9098
Epoch 30/80
188/188 [=====] - 5s 24ms/step - loss: 0.3084 - accuracy: 0.8713 - val_loss: 0.2998 - val_accuracy: 0.8888
Epoch 31/80
188/188 [=====] - 5s 24ms/step - loss: 0.3241 - accuracy: 0.8648 - val_loss: 0.3702 - val_accuracy: 0.8527
Epoch 32/80
188/188 [=====] - 5s 25ms/step - loss: 0.3175 - accuracy: 0.8705 - val_loss: 0.2866 - val_accuracy: 0.8938
Epoch 33/80
188/188 [=====] - 5s 25ms/step - loss: 0.3435 - accuracy: 0.8593 - val_loss: 0.4798 - val_accuracy: 0.8838
Epoch 34/80
188/188 [=====] - 5s 24ms/step - loss: 0.3433 - accuracy: 0.8573 - val_loss: 0.2443 - val_accuracy: 0.9048
Epoch 35/80
188/188 [=====] - 5s 25ms/step - loss: 0.3306 - accuracy: 0.8686 - val_loss: 2.3848 - val_accuracy: 0.7144
Epoch 36/80
188/188 [=====] - 5s 25ms/step - loss: 0.3549 - accuracy: 0.8630 - val_loss: 0.3681 - val_accuracy: 0.8798
Epoch 37/80
188/188 [=====] - 5s 25ms/step - loss: 0.3391 - accuracy: 0.8623 - val_loss: 0.3490 - val_accuracy: 0.8868
Epoch 38/80
188/188 [=====] - 5s 25ms/step - loss: 0.3542 - accuracy: 0.8596 - val_loss: 0.3876 - val_accuracy: 0.8397
Epoch 39/80
188/188 [=====] - 5s 25ms/step - loss: 0.3368 - accuracy: 0.8620 - val_loss: 0.5940 - val_accuracy: 0.8717
Epoch 40/80
188/188 [=====] - 5s 24ms/step - loss: 0.3310 - accuracy: 0.8706 - val_loss: 0.2680 - val_accuracy: 0.8958
Epoch 41/80
188/188 [=====] - 5s 24ms/step - loss: 0.3299 - accuracy: 0.8665 - val_loss: 0.3128 - val_accuracy: 0.8928
Epoch 42/80

188/188 [=====] - 5s 24ms/step - loss: 0.3455 - accuracy: 0.8573 - val_loss: 0.3010 - val_accuracy: 0.8918
Epoch 43/80
188/188 [=====] - 5s 24ms/step - loss: 0.3463 - accuracy: 0.8610 - val_loss: 0.4598 - val_accuracy: 0.8377
Epoch 44/80
188/188 [=====] - 5s 24ms/step - loss: 0.3580 - accuracy: 0.8603 - val_loss: 0.2928 - val_accuracy: 0.9018
Epoch 45/80
188/188 [=====] - 5s 25ms/step - loss: 0.3565 - accuracy: 0.8621 - val_loss: 0.3692 - val_accuracy: 0.8287
Epoch 46/80
188/188 [=====] - 5s 24ms/step - loss: 0.3547 - accuracy: 0.8610 - val_loss: 0.3488 - val_accuracy: 0.8798
Epoch 47/80
188/188 [=====] - 5s 24ms/step - loss: 0.3523 - accuracy: 0.8560 - val_loss: 0.5674 - val_accuracy: 0.7946
Epoch 48/80
188/188 [=====] - 5s 25ms/step - loss: 0.3737 - accuracy: 0.8553 - val_loss: 0.6356 - val_accuracy: 0.8146
Epoch 49/80
188/188 [=====] - 5s 24ms/step - loss: 0.3734 - accuracy: 0.8503 - val_loss: 0.9687 - val_accuracy: 0.7445
Epoch 50/80
188/188 [=====] - 5s 24ms/step - loss: 0.3659 - accuracy: 0.8525 - val_loss: 0.3908 - val_accuracy: 0.8427
Epoch 51/80
188/188 [=====] - 5s 25ms/step - loss: 0.3566 - accuracy: 0.8606 - val_loss: 0.9658 - val_accuracy: 0.7495
Epoch 52/80
188/188 [=====] - 5s 25ms/step - loss: 0.3662 - accuracy: 0.8521 - val_loss: 0.3714 - val_accuracy: 0.8968
Epoch 53/80
188/188 [=====] - 5s 24ms/step - loss: 0.3781 - accuracy: 0.8506 - val_loss: 0.2634 - val_accuracy: 0.9038
Epoch 54/80
188/188 [=====] - 5s 24ms/step - loss: 0.3798 - accuracy: 0.8456 - val_loss: 1.2284 - val_accuracy: 0.8176
Epoch 55/80
188/188 [=====] - 5s 25ms/step - loss: 0.4102 - accuracy: 0.8453 - val_loss: 0.2998 - val_accuracy: 0.9148
Epoch 56/80
188/188 [=====] - 5s 24ms/step - loss: 0.4404 - accuracy: 0.8251 - val_loss: 0.3784 - val_accuracy: 0.8828
Epoch 57/80
188/188 [=====] - 5s 24ms/step - loss: 0.3989 - accuracy: 0.8386 - val_loss: 0.2849 - val_accuracy: 0.8968
Epoch 58/80
188/188 [=====] - 5s 24ms/step - loss: 0.4147 - accuracy: 0.8403 - val_loss: 0.2906 - val_accuracy: 0.9058
Epoch 59/80
188/188 [=====] - 5s 25ms/step - loss: 0.3818 - accuracy: 0.8479 - val_loss: 0.4894 - val_accuracy: 0.8798

Epoch 60/80
188/188 [=====] - 5s 24ms/step - loss: 0.3876 - accuracy: 0.8358 - val_loss: 1.0239 - val_accuracy: 0.8267
Epoch 61/80
188/188 [=====] - 5s 25ms/step - loss: 0.4747 - accuracy: 0.8213 - val_loss: 0.3941 - val_accuracy: 0.8798
Epoch 62/80
188/188 [=====] - 5s 24ms/step - loss: 0.3936 - accuracy: 0.8324 - val_loss: 0.5380 - val_accuracy: 0.8257
Epoch 63/80
188/188 [=====] - 5s 25ms/step - loss: 0.4257 - accuracy: 0.8293 - val_loss: 0.4603 - val_accuracy: 0.8547
Epoch 64/80
188/188 [=====] - 5s 25ms/step - loss: 0.4307 - accuracy: 0.8311 - val_loss: 0.3651 - val_accuracy: 0.8206
Epoch 65/80
188/188 [=====] - 5s 25ms/step - loss: 0.4546 - accuracy: 0.8168 - val_loss: 0.2663 - val_accuracy: 0.8818
Epoch 66/80
188/188 [=====] - 5s 25ms/step - loss: 0.4361 - accuracy: 0.8286 - val_loss: 0.3358 - val_accuracy: 0.8868
Epoch 67/80
188/188 [=====] - 5s 24ms/step - loss: 0.4496 - accuracy: 0.8188 - val_loss: 0.3353 - val_accuracy: 0.8808
Epoch 68/80
188/188 [=====] - 5s 25ms/step - loss: 0.4687 - accuracy: 0.8269 - val_loss: 0.4740 - val_accuracy: 0.8667
Epoch 69/80
188/188 [=====] - 5s 24ms/step - loss: 0.4311 - accuracy: 0.8231 - val_loss: 0.4173 - val_accuracy: 0.8707
Epoch 70/80
188/188 [=====] - 5s 25ms/step - loss: 0.4982 - accuracy: 0.8074 - val_loss: 0.5288 - val_accuracy: 0.8236
Epoch 71/80
188/188 [=====] - 5s 24ms/step - loss: 0.4758 - accuracy: 0.8173 - val_loss: 0.4178 - val_accuracy: 0.8216
Epoch 72/80
188/188 [=====] - 5s 24ms/step - loss: 0.4796 - accuracy: 0.8183 - val_loss: 0.3006 - val_accuracy: 0.8888
Epoch 73/80
188/188 [=====] - 5s 25ms/step - loss: 0.4372 - accuracy: 0.8184 - val_loss: 0.4484 - val_accuracy: 0.8487
Epoch 74/80
188/188 [=====] - 5s 24ms/step - loss: 0.6609 - accuracy: 0.8118 - val_loss: 2.0788 - val_accuracy: 0.6703
Epoch 75/80
188/188 [=====] - 5s 25ms/step - loss: 0.5195 - accuracy: 0.8229 - val_loss: 0.7086 - val_accuracy: 0.8156
Epoch 76/80
188/188 [=====] - 5s 25ms/step - loss: 0.5815 - accuracy: 0.8069 - val_loss: 0.3858 - val_accuracy: 0.8858
Epoch 77/80

```
188/188 [=====] - 5s 25ms/step - loss: 0.5060 - accu
racy: 0.8174 - val_loss: 0.3619 - val_accuracy: 0.8367
Epoch 78/80
188/188 [=====] - 5s 24ms/step - loss: 0.4919 - accu
racy: 0.8014 - val_loss: 0.3146 - val_accuracy: 0.8727
Epoch 79/80
188/188 [=====] - 5s 25ms/step - loss: 0.4880 - accu
racy: 0.8144 - val_loss: 2.6971 - val_accuracy: 0.7094
Epoch 80/80
188/188 [=====] - 5s 24ms/step - loss: 0.4626 - accu
racy: 0.8159 - val_loss: 0.3114 - val_accuracy: 0.8657
```

Evaluating the model on the test set

In [20]:

```
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 14ms/step - loss: 0.3999 - accura
cy: 0.8868
Test accuracy: 0.887
```

Training a convnet from scratch with training sample of 5000, a validation sample of 500, and a test sample of 500

Downloading the data

```
!unzip -qq '/fs/ess/PGS0333/BA_64061_KSU/data/dogs-vs-cats.zip'
```

In [1]:

```
!unzip -qq train.zip
```

In [2]:

Copying images to training, validation, and test directories

```
import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index,
end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=5000)
make_subset("validation", start_index=5000, end_index=5500)
make_subset("test", start_index=5500, end_index=6000)
```

In [3]:

Building the model

Instantiating a small convnet for dogs vs. cats classification

```
from tensorflow import keras
from tensorflow.keras import layers

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
```

In [4]:

```
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

In [5]:

```
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 180, 180, 3)]	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 256)	0
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590080
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1)	12545
Total params: 991,041		
Trainable params: 991,041		
Non-trainable params: 0		

Configuring the model for training

In [6]:

```
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

Data preprocessing

Using image_dataset_from_directory to read images

In [7]:

```
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)

Found 10000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
```

In [8]:

```
import numpy as np
import tensorflow as tf
random_numbers = np.random.normal(size=(1000, 16))
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)
```

In [9]:

```
for i, element in enumerate(dataset):
    print(element.shape)
    if i >= 2:
        break

(16,)
(16,)
(16,)
```

In [10]:

```
batched_dataset = dataset.batch(32)
for i, element in enumerate(batched_dataset):
    print(element.shape)
    if i >= 2:
        break

(32, 16)
(32, 16)
(32, 16)
```

In [11]:

```
reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))
for i, element in enumerate(reshaped_dataset):
    print(element.shape)
    if i >= 2:
        break

(4, 4)
(4, 4)
(4, 4)
```

Displaying the shapes of the data and labels yielded by the Dataset

In [12]:

```
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    break
```

data batch shape: (32, 180, 180, 3)

labels batch shape: (32,)

Fitting the model using a Dataset

In [13]:

```
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)

Epoch 1/30
313/313 [=====] - 13s 25ms/step - loss: 0.7100 - accu
racy: 0.5490 - val_loss: 0.6479 - val_accuracy: 0.6480
Epoch 2/30
313/313 [=====] - 7s 22ms/step - loss: 0.6220 - accu
racy: 0.6644 - val_loss: 0.5241 - val_accuracy: 0.7440
Epoch 3/30
313/313 [=====] - 7s 22ms/step - loss: 0.5212 - accu
racy: 0.7444 - val_loss: 0.4661 - val_accuracy: 0.7810
Epoch 4/30
313/313 [=====] - 7s 22ms/step - loss: 0.4452 - accu
racy: 0.7956 - val_loss: 0.4746 - val_accuracy: 0.7870
Epoch 5/30
313/313 [=====] - 7s 23ms/step - loss: 0.3822 - accu
racy: 0.8326 - val_loss: 0.5290 - val_accuracy: 0.7740
Epoch 6/30
313/313 [=====] - 7s 22ms/step - loss: 0.3272 - accu
racy: 0.8573 - val_loss: 0.5761 - val_accuracy: 0.7860
Epoch 7/30
313/313 [=====] - 7s 22ms/step - loss: 0.2765 - accu
racy: 0.8835 - val_loss: 0.4929 - val_accuracy: 0.7930
Epoch 8/30
313/313 [=====] - 7s 22ms/step - loss: 0.2166 - accu
racy: 0.9101 - val_loss: 0.4669 - val_accuracy: 0.8280
Epoch 9/30
313/313 [=====] - 7s 22ms/step - loss: 0.1701 - accu
racy: 0.9297 - val_loss: 0.6306 - val_accuracy: 0.8150
Epoch 10/30
313/313 [=====] - 7s 22ms/step - loss: 0.1321 - accu
racy: 0.9482 - val_loss: 0.6038 - val_accuracy: 0.8430
Epoch 11/30
```


313/313 [=====] - 7s 23ms/step - loss: 0.1071 - accuracy: 0.9617 - val_loss: 0.6322 - val_accuracy: 0.8470
Epoch 12/30
313/313 [=====] - 7s 23ms/step - loss: 0.0902 - accuracy: 0.9680 - val_loss: 0.6687 - val_accuracy: 0.8510
Epoch 13/30
313/313 [=====] - 7s 23ms/step - loss: 0.0833 - accuracy: 0.9704 - val_loss: 0.7229 - val_accuracy: 0.8260
Epoch 14/30
313/313 [=====] - 7s 22ms/step - loss: 0.0759 - accuracy: 0.9729 - val_loss: 0.7790 - val_accuracy: 0.8110
Epoch 15/30
313/313 [=====] - 7s 22ms/step - loss: 0.0653 - accuracy: 0.9785 - val_loss: 0.7850 - val_accuracy: 0.8460
Epoch 16/30
313/313 [=====] - 7s 22ms/step - loss: 0.0643 - accuracy: 0.9800 - val_loss: 0.8671 - val_accuracy: 0.8410
Epoch 17/30
313/313 [=====] - 7s 22ms/step - loss: 0.0626 - accuracy: 0.9803 - val_loss: 1.5262 - val_accuracy: 0.8260
Epoch 18/30
313/313 [=====] - 7s 23ms/step - loss: 0.0674 - accuracy: 0.9806 - val_loss: 1.0057 - val_accuracy: 0.8490
Epoch 19/30
313/313 [=====] - 7s 22ms/step - loss: 0.0605 - accuracy: 0.9800 - val_loss: 1.0948 - val_accuracy: 0.8410
Epoch 20/30
313/313 [=====] - 7s 22ms/step - loss: 0.0641 - accuracy: 0.9825 - val_loss: 1.2610 - val_accuracy: 0.8370
Epoch 21/30
313/313 [=====] - 8s 24ms/step - loss: 0.0590 - accuracy: 0.9823 - val_loss: 1.3441 - val_accuracy: 0.8460
Epoch 22/30
313/313 [=====] - 7s 22ms/step - loss: 0.0557 - accuracy: 0.9862 - val_loss: 1.4184 - val_accuracy: 0.8470
Epoch 23/30
313/313 [=====] - 7s 22ms/step - loss: 0.0610 - accuracy: 0.9842 - val_loss: 1.3167 - val_accuracy: 0.8350
Epoch 24/30
313/313 [=====] - 7s 22ms/step - loss: 0.0581 - accuracy: 0.9847 - val_loss: 1.1253 - val_accuracy: 0.8530
Epoch 25/30
313/313 [=====] - 7s 23ms/step - loss: 0.0504 - accuracy: 0.9858 - val_loss: 1.3824 - val_accuracy: 0.8450
Epoch 26/30
313/313 [=====] - 7s 22ms/step - loss: 0.0646 - accuracy: 0.9854 - val_loss: 1.2968 - val_accuracy: 0.8640
Epoch 27/30
313/313 [=====] - 7s 22ms/step - loss: 0.0657 - accuracy: 0.9851 - val_loss: 1.4639 - val_accuracy: 0.8280
Epoch 28/30
313/313 [=====] - 7s 22ms/step - loss: 0.0579 - accuracy: 0.9862 - val_loss: 1.4657 - val_accuracy: 0.8490



```
Epoch 29/30
313/313 [=====] - 7s 22ms/step - loss: 0.0597 - accuracy: 0.9860 - val_loss: 1.5611 - val_accuracy: 0.8460
Epoch 30/30
313/313 [=====] - 7s 23ms/step - loss: 0.0515 - accuracy: 0.9861 - val_loss: 2.5216 - val_accuracy: 0.8380
```

Displaying curves of loss and accuracy during training

In [14]:

```
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```





Evaluating the model on the test set

In [15]:

```
test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 15ms/step - loss: 0.4459 - accuracy: 0.7990
Test accuracy: 0.799
```

Using data augmentation

Define a data augmentation stage to add to an image model

In [16]:

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
```

Displaying some randomly augmented training images

In [17]:

```
plt.figure(figsize=(10, 10))
for images, _ in train_dataset.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_images[0].numpy().astype("uint8"))
```

```
plt.axis("off")
```



Defining a new convnet that includes image augmentation, regularization and dropout

In [18]:

```
inputs = keras.Input(shape=(180, 180, 3))  
x = data_augmentation(inputs)  
x = layers.Rescaling(1./255)(x)
```

```

x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

Training the regularized convnet

In [19]:

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=80,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/80
313/313 [=====] - 8s 23ms/step - loss: 0.7085 - accuracy: 0.5472 - val_loss: 0.6473 - val_accuracy: 0.6280
Epoch 2/80
313/313 [=====] - 7s 23ms/step - loss: 0.6436 - accuracy: 0.6348 - val_loss: 0.5860 - val_accuracy: 0.6870
Epoch 3/80
313/313 [=====] - 7s 23ms/step - loss: 0.6026 - accuracy: 0.6773 - val_loss: 0.5413 - val_accuracy: 0.7280
Epoch 4/80
313/313 [=====] - 8s 25ms/step - loss: 0.5605 - accuracy: 0.7158 - val_loss: 0.6777 - val_accuracy: 0.6710
Epoch 5/80
313/313 [=====] - 7s 23ms/step - loss: 0.5249 - accuracy: 0.7430 - val_loss: 0.5795 - val_accuracy: 0.7050
Epoch 6/80
313/313 [=====] - 7s 23ms/step - loss: 0.4967 - accuracy: 0.7601 - val_loss: 0.4453 - val_accuracy: 0.7990
Epoch 7/80
313/313 [=====] - 7s 23ms/step - loss: 0.4714 - accuracy: 0.7775 - val_loss: 0.4263 - val_accuracy: 0.8020

Epoch 8/80
313/313 [=====] - 7s 23ms/step - loss: 0.4315 - accuracy: 0.8005 - val_loss: 0.3825 - val_accuracy: 0.8320
Epoch 9/80
313/313 [=====] - 7s 23ms/step - loss: 0.4126 - accuracy: 0.8136 - val_loss: 0.3821 - val_accuracy: 0.8300
Epoch 10/80
313/313 [=====] - 7s 23ms/step - loss: 0.3923 - accuracy: 0.8244 - val_loss: 0.9279 - val_accuracy: 0.6840
Epoch 11/80
313/313 [=====] - 7s 23ms/step - loss: 0.3709 - accuracy: 0.8406 - val_loss: 0.6279 - val_accuracy: 0.7800
Epoch 12/80
313/313 [=====] - 7s 23ms/step - loss: 0.3564 - accuracy: 0.8476 - val_loss: 0.5450 - val_accuracy: 0.8020
Epoch 13/80
313/313 [=====] - 7s 23ms/step - loss: 0.3380 - accuracy: 0.8495 - val_loss: 0.2726 - val_accuracy: 0.8870
Epoch 14/80
313/313 [=====] - 7s 23ms/step - loss: 0.3357 - accuracy: 0.8587 - val_loss: 0.2468 - val_accuracy: 0.9020
Epoch 15/80
313/313 [=====] - 7s 23ms/step - loss: 0.3135 - accuracy: 0.8667 - val_loss: 0.2766 - val_accuracy: 0.8790
Epoch 16/80
313/313 [=====] - 7s 23ms/step - loss: 0.3212 - accuracy: 0.8664 - val_loss: 1.1412 - val_accuracy: 0.6280
Epoch 17/80
313/313 [=====] - 7s 23ms/step - loss: 0.3134 - accuracy: 0.8695 - val_loss: 0.2261 - val_accuracy: 0.9020
Epoch 18/80
313/313 [=====] - 7s 23ms/step - loss: 0.3096 - accuracy: 0.8720 - val_loss: 0.2283 - val_accuracy: 0.9130
Epoch 19/80
313/313 [=====] - 7s 23ms/step - loss: 0.3088 - accuracy: 0.8674 - val_loss: 0.4922 - val_accuracy: 0.8490
Epoch 20/80
313/313 [=====] - 7s 23ms/step - loss: 0.3031 - accuracy: 0.8750 - val_loss: 0.2459 - val_accuracy: 0.8970
Epoch 21/80
313/313 [=====] - 7s 24ms/step - loss: 0.3020 - accuracy: 0.8723 - val_loss: 0.2751 - val_accuracy: 0.8980
Epoch 22/80
313/313 [=====] - 7s 23ms/step - loss: 0.3065 - accuracy: 0.8706 - val_loss: 0.2204 - val_accuracy: 0.9210
Epoch 23/80
313/313 [=====] - 7s 23ms/step - loss: 0.2973 - accuracy: 0.8759 - val_loss: 0.2631 - val_accuracy: 0.8910
Epoch 24/80
313/313 [=====] - 7s 23ms/step - loss: 0.3077 - accuracy: 0.8753 - val_loss: 0.2650 - val_accuracy: 0.8980
Epoch 25/80

313/313 [=====] - 7s 23ms/step - loss: 0.3067 - accuracy: 0.8758 - val_loss: 0.2240 - val_accuracy: 0.9170
Epoch 26/80
313/313 [=====] - 7s 23ms/step - loss: 0.3196 - accuracy: 0.8688 - val_loss: 0.2485 - val_accuracy: 0.8890
Epoch 27/80
313/313 [=====] - 7s 23ms/step - loss: 0.3124 - accuracy: 0.8727 - val_loss: 0.3176 - val_accuracy: 0.8990
Epoch 28/80
313/313 [=====] - 7s 23ms/step - loss: 0.3297 - accuracy: 0.8622 - val_loss: 0.2380 - val_accuracy: 0.8970
Epoch 29/80
313/313 [=====] - 7s 23ms/step - loss: 0.3271 - accuracy: 0.8665 - val_loss: 0.3226 - val_accuracy: 0.8960
Epoch 30/80
313/313 [=====] - 7s 23ms/step - loss: 0.3322 - accuracy: 0.8668 - val_loss: 0.4391 - val_accuracy: 0.8790
Epoch 31/80
313/313 [=====] - 7s 23ms/step - loss: 0.3331 - accuracy: 0.8623 - val_loss: 0.2398 - val_accuracy: 0.9170
Epoch 32/80
313/313 [=====] - 7s 23ms/step - loss: 0.3690 - accuracy: 0.8528 - val_loss: 0.3074 - val_accuracy: 0.8740
Epoch 33/80
313/313 [=====] - 8s 24ms/step - loss: 0.3580 - accuracy: 0.8530 - val_loss: 0.2445 - val_accuracy: 0.8970
Epoch 34/80
313/313 [=====] - 7s 23ms/step - loss: 0.3646 - accuracy: 0.8568 - val_loss: 0.3815 - val_accuracy: 0.8040
Epoch 35/80
313/313 [=====] - 7s 23ms/step - loss: 0.3949 - accuracy: 0.8419 - val_loss: 0.3172 - val_accuracy: 0.8670
Epoch 36/80
313/313 [=====] - 7s 23ms/step - loss: 0.3834 - accuracy: 0.8471 - val_loss: 0.2677 - val_accuracy: 0.8950
Epoch 37/80
313/313 [=====] - 8s 24ms/step - loss: 0.4062 - accuracy: 0.8369 - val_loss: 0.9175 - val_accuracy: 0.8380
Epoch 38/80
313/313 [=====] - 9s 27ms/step - loss: 0.3934 - accuracy: 0.8404 - val_loss: 0.3398 - val_accuracy: 0.9140
Epoch 39/80
313/313 [=====] - 9s 29ms/step - loss: 0.4115 - accuracy: 0.8388 - val_loss: 0.3155 - val_accuracy: 0.8780
Epoch 40/80
313/313 [=====] - 7s 23ms/step - loss: 0.4306 - accuracy: 0.8331 - val_loss: 0.6346 - val_accuracy: 0.7710
Epoch 41/80
313/313 [=====] - 7s 23ms/step - loss: 0.4223 - accuracy: 0.8352 - val_loss: 0.3256 - val_accuracy: 0.8680
Epoch 42/80
313/313 [=====] - 7s 23ms/step - loss: 0.4123 - accuracy: 0.8353 - val_loss: 0.2849 - val_accuracy: 0.8780

Epoch 43/80
313/313 [=====] - 7s 23ms/step - loss: 0.4292 - accuracy: 0.8374 - val_loss: 0.4640 - val_accuracy: 0.7860
Epoch 44/80
313/313 [=====] - 7s 23ms/step - loss: 0.4406 - accuracy: 0.8268 - val_loss: 0.3164 - val_accuracy: 0.8670
Epoch 45/80
313/313 [=====] - 7s 23ms/step - loss: 0.4382 - accuracy: 0.8288 - val_loss: 0.6347 - val_accuracy: 0.8510
Epoch 46/80
313/313 [=====] - 7s 23ms/step - loss: 0.4506 - accuracy: 0.8228 - val_loss: 0.4618 - val_accuracy: 0.8030
Epoch 47/80
313/313 [=====] - 7s 23ms/step - loss: 0.4790 - accuracy: 0.8111 - val_loss: 0.3703 - val_accuracy: 0.8700
Epoch 48/80
313/313 [=====] - 8s 24ms/step - loss: 0.4737 - accuracy: 0.8153 - val_loss: 0.4769 - val_accuracy: 0.8420
Epoch 49/80
313/313 [=====] - 8s 24ms/step - loss: 0.5042 - accuracy: 0.8108 - val_loss: 0.4332 - val_accuracy: 0.8170
Epoch 50/80
313/313 [=====] - 8s 24ms/step - loss: 0.4661 - accuracy: 0.8106 - val_loss: 0.3641 - val_accuracy: 0.8390
Epoch 51/80
313/313 [=====] - 7s 23ms/step - loss: 0.5111 - accuracy: 0.8038 - val_loss: 0.5490 - val_accuracy: 0.7600
Epoch 52/80
313/313 [=====] - 7s 23ms/step - loss: 0.4933 - accuracy: 0.8100 - val_loss: 0.4601 - val_accuracy: 0.8140
Epoch 53/80
313/313 [=====] - 7s 23ms/step - loss: 0.5665 - accuracy: 0.7935 - val_loss: 0.5204 - val_accuracy: 0.8180
Epoch 54/80
313/313 [=====] - 8s 24ms/step - loss: 0.8305 - accuracy: 0.7883 - val_loss: 0.4358 - val_accuracy: 0.8030
Epoch 55/80
313/313 [=====] - 8s 24ms/step - loss: 0.5653 - accuracy: 0.7699 - val_loss: 1.1045 - val_accuracy: 0.7030
Epoch 56/80
313/313 [=====] - 7s 23ms/step - loss: 0.5717 - accuracy: 0.7654 - val_loss: 0.6066 - val_accuracy: 0.6160
Epoch 57/80
313/313 [=====] - 7s 23ms/step - loss: 0.6231 - accuracy: 0.7561 - val_loss: 0.7035 - val_accuracy: 0.8050
Epoch 58/80
313/313 [=====] - 7s 23ms/step - loss: 0.6366 - accuracy: 0.7521 - val_loss: 0.3822 - val_accuracy: 0.8320
Epoch 59/80
313/313 [=====] - 7s 23ms/step - loss: 0.7207 - accuracy: 0.7526 - val_loss: 0.4978 - val_accuracy: 0.7790
Epoch 60/80

313/313 [=====] - 7s 23ms/step - loss: 0.5988 - accuracy: 0.7567 - val_loss: 0.7791 - val_accuracy: 0.8630
Epoch 61/80
313/313 [=====] - 8s 24ms/step - loss: 0.7173 - accuracy: 0.7496 - val_loss: 0.3506 - val_accuracy: 0.8240
Epoch 62/80
313/313 [=====] - 7s 23ms/step - loss: 0.6413 - accuracy: 0.7624 - val_loss: 0.5071 - val_accuracy: 0.7890
Epoch 63/80
313/313 [=====] - 7s 23ms/step - loss: 0.6937 - accuracy: 0.7424 - val_loss: 0.5318 - val_accuracy: 0.7410
Epoch 64/80
313/313 [=====] - 7s 23ms/step - loss: 0.5900 - accuracy: 0.7461 - val_loss: 0.4062 - val_accuracy: 0.8290
Epoch 65/80
313/313 [=====] - 7s 23ms/step - loss: 0.6657 - accuracy: 0.7406 - val_loss: 0.4436 - val_accuracy: 0.7760
Epoch 66/80
313/313 [=====] - 7s 24ms/step - loss: 0.6120 - accuracy: 0.7284 - val_loss: 0.3400 - val_accuracy: 0.8570
Epoch 67/80
313/313 [=====] - 7s 23ms/step - loss: 0.6894 - accuracy: 0.7491 - val_loss: 0.5961 - val_accuracy: 0.7390
Epoch 68/80
313/313 [=====] - 7s 23ms/step - loss: 0.6624 - accuracy: 0.7276 - val_loss: 0.4832 - val_accuracy: 0.7910
Epoch 69/80
313/313 [=====] - 7s 23ms/step - loss: 0.6145 - accuracy: 0.7428 - val_loss: 0.4742 - val_accuracy: 0.7550
Epoch 70/80
313/313 [=====] - 8s 26ms/step - loss: 0.7200 - accuracy: 0.7363 - val_loss: 0.5561 - val_accuracy: 0.7350
Epoch 71/80
313/313 [=====] - 8s 24ms/step - loss: 0.6692 - accuracy: 0.7420 - val_loss: 0.5122 - val_accuracy: 0.7760
Epoch 72/80
313/313 [=====] - 8s 24ms/step - loss: 0.6900 - accuracy: 0.7354 - val_loss: 2.8342 - val_accuracy: 0.5990
Epoch 73/80
313/313 [=====] - 8s 24ms/step - loss: 0.6887 - accuracy: 0.7213 - val_loss: 0.5601 - val_accuracy: 0.7590
Epoch 74/80
313/313 [=====] - 7s 23ms/step - loss: 0.6664 - accuracy: 0.7248 - val_loss: 1.0051 - val_accuracy: 0.5290
Epoch 75/80
313/313 [=====] - 8s 24ms/step - loss: 0.6987 - accuracy: 0.7268 - val_loss: 0.5395 - val_accuracy: 0.7330
Epoch 76/80
313/313 [=====] - 8s 24ms/step - loss: 0.7089 - accuracy: 0.7017 - val_loss: 0.5606 - val_accuracy: 0.6710
Epoch 77/80
313/313 [=====] - 8s 26ms/step - loss: 0.6801 - accuracy: 0.7002 - val_loss: 0.4867 - val_accuracy: 0.7900


```
Epoch 78/80
313/313 [=====] - 9s 27ms/step - loss: 0.6834 - accuracy: 0.7238 - val_loss: 0.6524 - val_accuracy: 0.7180
Epoch 79/80
313/313 [=====] - 8s 24ms/step - loss: 0.7559 - accuracy: 0.7261 - val_loss: 0.5264 - val_accuracy: 0.7770
Epoch 80/80
313/313 [=====] - 8s 24ms/step - loss: 0.6554 - accuracy: 0.7236 - val_loss: 0.5299 - val_accuracy: 0.7480
```

Evaluating the model on the test set

In [20]:

```
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 15ms/step - loss: 0.2071 - accuracy: 0.9150
Test accuracy: 0.915
```

Training a convnet from scratch with training sample of 8000, a validation sample of 500, and a test sample of 500

Downloading the data

```
#!/unzip -qq '/fs/ess/PGS0333/BA_64061_KSU/data/dogs-vs-cats.zip'
```

In [1]:

```
#!/unzip -qq train.zip
```

In [2]:

Copying images to training, validation, and test directories

```
import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index,
end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=7999)
make_subset("validation", start_index=8000, end_index=8499)
make_subset("test", start_index=8499, end_index=8999)
```

In [3]:

Building the model

Instantiating a small convnet for dogs vs. cats classification

```
from tensorflow import keras
from tensorflow.keras import layers

inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
```

In [4]:

```
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

In [5]:

```
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 180, 180, 3)]	0
rescaling (Rescaling)	(None, 180, 180, 3)	0
conv2d (Conv2D)	(None, 178, 178, 32)	896
max_pooling2d (MaxPooling2D)	(None, 89, 89, 32)	0
conv2d_1 (Conv2D)	(None, 87, 87, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 43, 43, 64)	0
conv2d_2 (Conv2D)	(None, 41, 41, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 20, 20, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 256)	0
conv2d_4 (Conv2D)	(None, 7, 7, 256)	590080
flatten (Flatten)	(None, 12544)	0
dense (Dense)	(None, 1)	12545
=====		
Total params: 991,041		
Trainable params: 991,041		
Non-trainable params: 0		

Configuring the model for training

In [6]:

```
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
```

Data preprocessing

Using image_dataset_from_directory to read images

In [7]:

```
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",
    image_size=(180, 180),
    batch_size=32)

Found 15998 files belonging to 2 classes.
Found 998 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
```

In [8]:

```
import numpy as np
import tensorflow as tf
random_numbers = np.random.normal(size=(1000, 16))
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)
```

In [9]:

```
for i, element in enumerate(dataset):
    print(element.shape)
    if i >= 2:
        break

(16,)
(16,)
(16,)
```

In [10]:

```
batched_dataset = dataset.batch(32)
for i, element in enumerate(batched_dataset):
    print(element.shape)
    if i >= 2:
        break

(32, 16)
(32, 16)
(32, 16)
```

In [11]:

```
reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))
for i, element in enumerate(reshaped_dataset):
    print(element.shape)
    if i >= 2:
        break

(4, 4)
(4, 4)
(4, 4)
```

Displaying the shapes of the data and labels yielded by the Dataset

In [12]:

```
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    break
```

data batch shape: (32, 180, 180, 3)

labels batch shape: (32,)

Fitting the model using a Dataset

In [13]:

```
"""
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)
"""
```

Out[13]:

```
'\ncallbacks = [\n    keras.callbacks.ModelCheckpoint(\n        filepath="convnet_from_scratch.keras",\n        save_best_only=True,\n        monitor="val_loss")\n]\nhistory = model.fit(\n    train_dataset,\n    epochs=30,\n    validation_data=validation_dataset,\n    callbacks=callbacks)\n'
```

Displaying curves of loss and accuracy during training

In [14]:

```
"""
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
"""
```

Out[14]:

```
'\nimport matplotlib.pyplot as plt\naccuracy = history.history["accuracy"]\nval_accuracy = history.history["val_accuracy"]\nloss = history.history["loss"]
```

```
\nval_loss = history.history["val_loss"]\nepochs = range(1, len(accuracy) + 1)\n\nplt.plot(epochs, accuracy, "bo", label="Training accuracy")\nplt.plot(epochs, val_accuracy, "b", label="Validation accuracy")\nplt.title("Training and validation accuracy")\nplt.legend()\nplt.figure()\nplt.plot(epochs, loss, "bo", label="Training loss")\nplt.plot(epochs, val_loss, "b", label="Validation loss")\nplt.title("Training and validation loss")\nplt.legend()\nplt.show()\n'
```

Evaluating the model on the test set

In [15]:

```
"""\ntest_model = keras.models.load_model("convnet_from_scratch.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n"""
```

Out[15]:

```
'\ntest_model = keras.models.load_model("convnet_from_scratch.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n'
```

Using data augmentation

Define a data augmentation stage to add to an image model

In [16]:

```
data_augmentation = keras.Sequential(\n    [\n        layers.RandomFlip("horizontal"),\n        layers.RandomRotation(0.1),\n        layers.RandomZoom(0.2),\n    ]\n)
```

Displaying some randomly augmented training images

In [17]:

```
"""\nplt.figure(figsize=(10, 10))\nfor images, _ in train_dataset.take(1):\n    for i in range(9):\n        augmented_images = data_augmentation(images)\n        ax = plt.subplot(3, 3, i + 1)\n        plt.imshow(augmented_images[0].numpy().astype("uint8"))\n        plt.axis("off")\n"""
```

Out[17]:

```
'\nplt.figure(figsize=(10, 10))\nfor images, _ in train_dataset.take(1):\n    for i in range(9):\n        augmented_images = data_augmentation(images)\n        ax = plt.subplot(3, 3, i + 1)\n        plt.imshow(augmented_images[0].numpy().astype("uint8"))\n        plt.axis("off")\n'
```

Defining a new convnet that includes image augmentation, regularization and dropout

In [18]:

```
inputs = keras.Input(shape=(180, 180, 3))
```

```

x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

Training the regularized convnet

In [19]:

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=80,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/80
500/500 [=====] - 16s 24ms/step - loss: 0.6698 - accuracy: 0.5857 - val_loss: 0.5948 - val_accuracy: 0.6974
Epoch 2/80
500/500 [=====] - 12s 23ms/step - loss: 0.5821 - accuracy: 0.6964 - val_loss: 0.5041 - val_accuracy: 0.7465
Epoch 3/80
500/500 [=====] - 11s 23ms/step - loss: 0.5173 - accuracy: 0.7536 - val_loss: 0.4001 - val_accuracy: 0.8196
Epoch 4/80
500/500 [=====] - 12s 23ms/step - loss: 0.4723 - accuracy: 0.7782 - val_loss: 0.4987 - val_accuracy: 0.7836
Epoch 5/80
500/500 [=====] - 12s 23ms/step - loss: 0.4270 - accuracy: 0.8062 - val_loss: 0.3770 - val_accuracy: 0.8367
Epoch 6/80
500/500 [=====] - 12s 23ms/step - loss: 0.3997 - accuracy: 0.8220 - val_loss: 0.2858 - val_accuracy: 0.8818
Epoch 7/80

500/500 [=====] - 12s 23ms/step - loss: 0.3660 - accuracy: 0.8396 - val_loss: 0.3034 - val_accuracy: 0.8758
Epoch 8/80
500/500 [=====] - 11s 22ms/step - loss: 0.3541 - accuracy: 0.8462 - val_loss: 0.2472 - val_accuracy: 0.8988
Epoch 9/80
500/500 [=====] - 12s 23ms/step - loss: 0.3404 - accuracy: 0.8548 - val_loss: 0.3367 - val_accuracy: 0.8567
Epoch 10/80
500/500 [=====] - 12s 23ms/step - loss: 0.3285 - accuracy: 0.8629 - val_loss: 0.2327 - val_accuracy: 0.8998
Epoch 11/80
500/500 [=====] - 11s 23ms/step - loss: 0.3237 - accuracy: 0.8645 - val_loss: 0.2286 - val_accuracy: 0.9058
Epoch 12/80
500/500 [=====] - 12s 23ms/step - loss: 0.3222 - accuracy: 0.8661 - val_loss: 0.2522 - val_accuracy: 0.9038
Epoch 13/80
500/500 [=====] - 12s 23ms/step - loss: 0.3229 - accuracy: 0.8694 - val_loss: 0.2510 - val_accuracy: 0.9128
Epoch 14/80
500/500 [=====] - 12s 23ms/step - loss: 0.3242 - accuracy: 0.8648 - val_loss: 0.2271 - val_accuracy: 0.9098
Epoch 15/80
500/500 [=====] - 11s 21ms/step - loss: 0.3219 - accuracy: 0.8669 - val_loss: 0.4112 - val_accuracy: 0.8246
Epoch 16/80
500/500 [=====] - 10s 21ms/step - loss: 0.3310 - accuracy: 0.8598 - val_loss: 0.1919 - val_accuracy: 0.9228
Epoch 17/80
500/500 [=====] - 11s 22ms/step - loss: 0.3342 - accuracy: 0.8607 - val_loss: 0.2056 - val_accuracy: 0.9188
Epoch 18/80
500/500 [=====] - 10s 21ms/step - loss: 0.3450 - accuracy: 0.8559 - val_loss: 0.2414 - val_accuracy: 0.8908
Epoch 19/80
500/500 [=====] - 10s 21ms/step - loss: 0.3657 - accuracy: 0.8515 - val_loss: 0.3468 - val_accuracy: 0.8307
Epoch 20/80
500/500 [=====] - 10s 21ms/step - loss: 0.3517 - accuracy: 0.8532 - val_loss: 0.2270 - val_accuracy: 0.9208
Epoch 21/80
500/500 [=====] - 11s 21ms/step - loss: 0.3681 - accuracy: 0.8481 - val_loss: 0.3836 - val_accuracy: 0.8888
Epoch 22/80
500/500 [=====] - 10s 21ms/step - loss: 0.3722 - accuracy: 0.8447 - val_loss: 0.6300 - val_accuracy: 0.7285
Epoch 23/80
500/500 [=====] - 10s 20ms/step - loss: 0.4038 - accuracy: 0.8415 - val_loss: 0.5272 - val_accuracy: 0.8717
Epoch 24/80
500/500 [=====] - 10s 21ms/step - loss: 0.4020 - accuracy: 0.8377 - val_loss: 0.2686 - val_accuracy: 0.8988

Epoch 25/80
500/500 [=====] - 10s 21ms/step - loss: 0.4430 - accuracy: 0.8336 - val_loss: 1.5721 - val_accuracy: 0.7806
Epoch 26/80
500/500 [=====] - 10s 21ms/step - loss: 0.4454 - accuracy: 0.8182 - val_loss: 0.3221 - val_accuracy: 0.8707
Epoch 27/80
500/500 [=====] - 11s 21ms/step - loss: 0.4203 - accuracy: 0.8275 - val_loss: 0.2136 - val_accuracy: 0.9118
Epoch 28/80
500/500 [=====] - 11s 21ms/step - loss: 0.4422 - accuracy: 0.8264 - val_loss: 0.2804 - val_accuracy: 0.8788
Epoch 29/80
500/500 [=====] - 10s 21ms/step - loss: 0.4528 - accuracy: 0.8213 - val_loss: 0.3431 - val_accuracy: 0.8677
Epoch 30/80
500/500 [=====] - 11s 21ms/step - loss: 0.4511 - accuracy: 0.8179 - val_loss: 0.3673 - val_accuracy: 0.8878
Epoch 31/80
500/500 [=====] - 10s 20ms/step - loss: 0.4706 - accuracy: 0.8105 - val_loss: 0.4326 - val_accuracy: 0.8427
Epoch 32/80
500/500 [=====] - 10s 20ms/step - loss: 0.4646 - accuracy: 0.8095 - val_loss: 0.3251 - val_accuracy: 0.8707
Epoch 33/80
500/500 [=====] - 10s 21ms/step - loss: 0.5123 - accuracy: 0.8054 - val_loss: 1.0964 - val_accuracy: 0.6393
Epoch 34/80
500/500 [=====] - 10s 20ms/step - loss: 0.5418 - accuracy: 0.7726 - val_loss: 0.3325 - val_accuracy: 0.8497
Epoch 35/80
500/500 [=====] - 10s 20ms/step - loss: 0.4987 - accuracy: 0.8029 - val_loss: 0.3348 - val_accuracy: 0.8617
Epoch 36/80
500/500 [=====] - 11s 21ms/step - loss: 0.5345 - accuracy: 0.8021 - val_loss: 0.4660 - val_accuracy: 0.8136
Epoch 37/80
500/500 [=====] - 10s 21ms/step - loss: 0.5108 - accuracy: 0.7906 - val_loss: 0.4385 - val_accuracy: 0.8126
Epoch 38/80
500/500 [=====] - 10s 21ms/step - loss: 0.5487 - accuracy: 0.7854 - val_loss: 0.3410 - val_accuracy: 0.8507
Epoch 39/80
500/500 [=====] - 10s 21ms/step - loss: 0.5248 - accuracy: 0.7893 - val_loss: 0.3030 - val_accuracy: 0.8697
Epoch 40/80
500/500 [=====] - 10s 21ms/step - loss: 0.5730 - accuracy: 0.7750 - val_loss: 0.3053 - val_accuracy: 0.8687
Epoch 41/80
500/500 [=====] - 10s 21ms/step - loss: 0.5474 - accuracy: 0.7735 - val_loss: 0.4701 - val_accuracy: 0.7595
Epoch 42/80

500/500 [=====] - 10s 20ms/step - loss: 0.5688 - accuracy: 0.7653 - val_loss: 0.5299 - val_accuracy: 0.8327
Epoch 43/80
500/500 [=====] - 10s 20ms/step - loss: 0.5633 - accuracy: 0.7732 - val_loss: 0.5697 - val_accuracy: 0.7154
Epoch 44/80
500/500 [=====] - 10s 20ms/step - loss: 0.6075 - accuracy: 0.7611 - val_loss: 0.3907 - val_accuracy: 0.8267
Epoch 45/80
500/500 [=====] - 10s 20ms/step - loss: 0.6059 - accuracy: 0.7528 - val_loss: 0.4196 - val_accuracy: 0.8317
Epoch 46/80
500/500 [=====] - 10s 20ms/step - loss: 0.7011 - accuracy: 0.7299 - val_loss: 0.4721 - val_accuracy: 0.8006
Epoch 47/80
500/500 [=====] - 11s 21ms/step - loss: 0.7024 - accuracy: 0.7442 - val_loss: 0.4047 - val_accuracy: 0.8086
Epoch 48/80
500/500 [=====] - 10s 21ms/step - loss: 0.7302 - accuracy: 0.7452 - val_loss: 0.5936 - val_accuracy: 0.7004
Epoch 49/80
500/500 [=====] - 10s 20ms/step - loss: 0.7334 - accuracy: 0.7185 - val_loss: 0.3979 - val_accuracy: 0.8347
Epoch 50/80
500/500 [=====] - 10s 21ms/step - loss: 0.6484 - accuracy: 0.7150 - val_loss: 0.4708 - val_accuracy: 0.8046
Epoch 51/80
500/500 [=====] - 10s 20ms/step - loss: 0.6544 - accuracy: 0.7255 - val_loss: 0.4457 - val_accuracy: 0.8036
Epoch 52/80
500/500 [=====] - 10s 20ms/step - loss: 0.6426 - accuracy: 0.7240 - val_loss: 0.4706 - val_accuracy: 0.8196
Epoch 53/80
500/500 [=====] - 10s 21ms/step - loss: 0.6671 - accuracy: 0.7307 - val_loss: 0.5251 - val_accuracy: 0.7926
Epoch 54/80
500/500 [=====] - 10s 21ms/step - loss: 0.6281 - accuracy: 0.7216 - val_loss: 0.4577 - val_accuracy: 0.7966
Epoch 55/80
500/500 [=====] - 10s 21ms/step - loss: 0.7185 - accuracy: 0.6918 - val_loss: 0.5353 - val_accuracy: 0.7405
Epoch 56/80
500/500 [=====] - 10s 21ms/step - loss: 0.7301 - accuracy: 0.7176 - val_loss: 0.4469 - val_accuracy: 0.7866
Epoch 57/80
500/500 [=====] - 10s 20ms/step - loss: 0.6787 - accuracy: 0.6987 - val_loss: 0.7591 - val_accuracy: 0.8377
Epoch 58/80
500/500 [=====] - 10s 21ms/step - loss: 0.8962 - accuracy: 0.6985 - val_loss: 0.4884 - val_accuracy: 0.7936
Epoch 59/80
500/500 [=====] - 10s 20ms/step - loss: 0.8222 - accuracy: 0.6948 - val_loss: 0.5391 - val_accuracy: 0.8026

Epoch 60/80
500/500 [=====] - 10s 20ms/step - loss: 0.7255 - accuracy: 0.6885 - val_loss: 0.5430 - val_accuracy: 0.7315
Epoch 61/80
500/500 [=====] - 10s 21ms/step - loss: 0.7795 - accuracy: 0.6950 - val_loss: 0.4455 - val_accuracy: 0.8056
Epoch 62/80
500/500 [=====] - 11s 21ms/step - loss: 0.7661 - accuracy: 0.6893 - val_loss: 0.4932 - val_accuracy: 0.7996
Epoch 63/80
500/500 [=====] - 10s 21ms/step - loss: 0.6789 - accuracy: 0.6858 - val_loss: 0.7269 - val_accuracy: 0.7325
Epoch 64/80
500/500 [=====] - 10s 21ms/step - loss: 0.6541 - accuracy: 0.7032 - val_loss: 0.4906 - val_accuracy: 0.7575
Epoch 65/80
500/500 [=====] - 10s 20ms/step - loss: 0.7988 - accuracy: 0.7289 - val_loss: 0.4875 - val_accuracy: 0.7766
Epoch 66/80
500/500 [=====] - 10s 20ms/step - loss: 0.7619 - accuracy: 0.7186 - val_loss: 0.5024 - val_accuracy: 0.7856
Epoch 67/80
500/500 [=====] - 10s 21ms/step - loss: 0.6605 - accuracy: 0.7142 - val_loss: 0.5712 - val_accuracy: 0.8126
Epoch 68/80
500/500 [=====] - 10s 21ms/step - loss: 0.7300 - accuracy: 0.6925 - val_loss: 0.8057 - val_accuracy: 0.8146
Epoch 69/80
500/500 [=====] - 10s 20ms/step - loss: 0.7359 - accuracy: 0.7048 - val_loss: 0.4732 - val_accuracy: 0.7826
Epoch 70/80
500/500 [=====] - 11s 21ms/step - loss: 0.6933 - accuracy: 0.7053 - val_loss: 0.5170 - val_accuracy: 0.7485
Epoch 71/80
500/500 [=====] - 10s 20ms/step - loss: 0.8406 - accuracy: 0.7012 - val_loss: 0.9676 - val_accuracy: 0.6253
Epoch 72/80
500/500 [=====] - 10s 20ms/step - loss: 0.6573 - accuracy: 0.7063 - val_loss: 1.1291 - val_accuracy: 0.6914
Epoch 73/80
500/500 [=====] - 10s 20ms/step - loss: 0.8186 - accuracy: 0.6640 - val_loss: 0.6766 - val_accuracy: 0.5561
Epoch 74/80
500/500 [=====] - 11s 21ms/step - loss: 0.9503 - accuracy: 0.6762 - val_loss: 0.5581 - val_accuracy: 0.7375
Epoch 75/80
500/500 [=====] - 10s 21ms/step - loss: 0.7765 - accuracy: 0.6905 - val_loss: 0.7470 - val_accuracy: 0.6443
Epoch 76/80
500/500 [=====] - 10s 21ms/step - loss: 0.7347 - accuracy: 0.6849 - val_loss: 0.4996 - val_accuracy: 0.8126
Epoch 77/80

```
500/500 [=====] - 10s 20ms/step - loss: 0.7122 - acc
uracy: 0.6692 - val_loss: 0.9294 - val_accuracy: 0.7615
Epoch 78/80
500/500 [=====] - 11s 22ms/step - loss: 0.7860 - acc
uracy: 0.6724 - val_loss: 0.5466 - val_accuracy: 0.7365
Epoch 79/80
500/500 [=====] - 10s 21ms/step - loss: 0.9206 - acc
uracy: 0.6649 - val_loss: 0.5738 - val_accuracy: 0.6834
Epoch 80/80
500/500 [=====] - 10s 20ms/step - loss: 0.8316 - acc
uracy: 0.6814 - val_loss: 0.5858 - val_accuracy: 0.7555
```

Evaluating the model on the test set

In [20]:

```
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 15ms/step - loss: 0.2239 - accura
cy: 0.9070
Test accuracy: 0.907
```

Using a pretrained network with training sample of 1000, a validation sample of 500, and a test sample of 500

Downloading the data

```
#!/unzip -qq '/fs/ess/PGS0333/BA_64061_KSU/data/dogs-vs-cats.zip'
```

In [1]:

```
#!/unzip -qq train.zip
```

In [2]:

Copying images to training, validation, and test directories

```
import os, shutil, pathlib
```

In [3]:

```
original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index,
end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=1000)
make_subset("validation", start_index=1000, end_index=1500)
make_subset("test", start_index=1500, end_index=2000)
```

Building the model

Instantiating a small convnet for dogs vs. cats classification

```
"""
from tensorflow import keras
from tensorflow.keras import layers
inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
```

In [4]:

```

x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
"""

```

Out[4]:

```

'\nfrom tensorflow import keras\nfrom tensorflow.keras import layers\ninputs
= keras.Input(shape=(180, 180, 3))\nx = layers.Rescaling(1./255)(inputs)\nx =
layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)\nx = layers.M
axPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filters=64, kernel_size=3, act
ivation="relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2
D(filters=128, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling2D
(pool_size=2)(x)\nx = layers.Conv2D(filters=256, kernel_size=3, activation="r
elu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filters=
256, kernel_size=3, activation="relu")(x)\nx = layers.Flatten()(x)\noutputs =
layers.Dense(1, activation="sigmoid")(x)\nmodel = keras.Model(inputs=inputs,
outputs=outputs)\n'

```

In [5]:

```

# model.summary()

```

Configuring the model for training

In [6]:

```

"""
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
"""

```

Out[6]:

```

'\nmodel.compile(loss="binary_crossentropy",\n
optimizer="rmspro
p",\n
metrics=["accuracy"])\n'

```

Data preprocessing

Using image_dataset_from_directory to read images

In [7]:

```

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",

```

```

        image_size=(180, 180),
        batch_size=32)

```

```

Found 2000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.
Found 1000 files belonging to 2 classes.

```

In [8]:

```

"""
import numpy as np
import tensorflow as tf
random_numbers = np.random.normal(size=(1000, 16))
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)
"""

```

Out[8]:

```

'\nimport numpy as np\nimport tensorflow as tf\nrandom_numbers = np.random.no
rmal(size=(1000, 16))\ndataset = tf.data.Dataset.from_tensor_slices(random_nu
mbers)\n'

```

In [9]:

```

"""
for i, element in enumerate(dataset):
    print(element.shape)
    if i >= 2:
        break
"""

```

Out[9]:

```

'\nfor i, element in enumerate(dataset):\n    print(element.shape)\n    if i >
= 2:\n        break\n'

```

In [10]:

```

"""
batched_dataset = dataset.batch(32)
for i, element in enumerate(batched_dataset):
    print(element.shape)
    if i >= 2:
        break
"""

```

Out[10]:

```

'\nbatched_dataset = dataset.batch(32)\nfor i, element in enumerate(batched_d
ataset):\n    print(element.shape)\n    if i >= 2:\n        break\n'

```

In [11]:

```

"""
reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))
for i, element in enumerate(reshaped_dataset):
    print(element.shape)
    if i >= 2:
        break
"""

```

Out[11]:

```

'\nreshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))\nfor i, el
ement in enumerate(reshaped_dataset):\n    print(element.shape)\n    if i >=
2:\n        break\n'

```

Displaying the shapes of the data and labels yielded by the Dataset

In [12]:

```
"""
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    break
"""
```

Out[12]:

```
'\nfor data_batch, labels_batch in train_dataset:\n    print("data batch shape:", data_batch.shape)\n    print("labels batch shape:", labels_batch.shape)\n    break\n'
```

Fitting the model using a Dataset

In [13]:

```
"""
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)
"""
```

Out[13]:

```
'\ncallbacks = [\n    keras.callbacks.ModelCheckpoint(\n        filepath="convnet_from_scratch.keras",\n        save_best_only=True,\n        monitor="val_loss")\n]\nhistory = model.fit(\n    train_dataset,\n    epochs=30,\n    validation_data=validation_dataset,\n    callbacks=callbacks)\n'
```

Displaying curves of loss and accuracy during training

In [14]:

```
"""
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
"""
```



```
plt.show()
"""
```

Out[14]:

```
'\nimport matplotlib.pyplot as plt\naccuracy = history.history["accuracy"]\nval_accuracy = history.history["val_accuracy"]\nloss = history.history["loss"]\nval_loss = history.history["val_loss"]\nepochs = range(1, len(accuracy) + 1)\nplt.plot(epochs, accuracy, "bo", label="Training accuracy")\nplt.plot(epochs, val_accuracy, "b", label="Validation accuracy")\nplt.title("Training and validation accuracy")\nplt.legend()\nplt.figure()\nplt.plot(epochs, loss, "bo", label="Training loss")\nplt.plot(epochs, val_loss, "b", label="Validation loss")\nplt.title("Training and validation loss")\nplt.legend()\nplt.show()\n'
```

Evaluating the model on the test set

In [15]:

```
"""
test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
"""
```

Out[15]:

```
'\ntest_model = keras.models.load_model("convnet_from_scratch.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n'
```

Using data augmentation

Define a data augmentation stage to add to an image model

In [16]:

```
"""
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
"""
```

Out[16]:

```
'\ndata_augmentation = keras.Sequential(\n    [\n        layers.RandomFlip("horizontal"),\n        layers.RandomRotation(0.1),\n        layers.RandomZoom(0.2),\n    ]\n)\n'
```

Displaying some randomly augmented training images

In [17]:

```
"""
plt.figure(figsize=(10, 10))
for images, _ in train_dataset.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
```

```

plt.imshow(augmented_images[0].numpy().astype("uint8"))
plt.axis("off")
"""

```

Out[17]:

```

'\nplt.figure(figsize=(10, 10))\nfor images, _ in train_dataset.take(1):\n
  for i in range(9):\n      augmented_images = data_augmentation(images)\n
      ax = plt.subplot(3, 3, i + 1)\n      plt.imshow(augmented_images[0].n
umpy().astype("uint8"))\n      plt.axis("off")\n'

```

Defining a new convnet that includes image augmentation and dropout

In [18]:

```

"""
inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
"""

```

Out[18]:

```

'\ninputs = keras.Input(shape=(180, 180, 3))\nx = data_augmentation(inputs)\n
x = layers.Rescaling(1./255)(x)\nx = layers.Conv2D(filters=32, kernel_size=3,
  activation="relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.C
onv2D(filters=64, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling
2D(pool_size=2)(x)\nx = layers.Conv2D(filters=128, kernel_size=3, activation=
"relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filter
s=256, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling2D(pool_siz
e=2)(x)\nx = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)\
nx = layers.Flatten()(x)\nx = layers.Dropout(0.5)(x)\noutputs = layers.Dense
(1, activation="sigmoid")(x)\nmodel = keras.Model(inputs=inputs, outputs=outp
uts)\n\nmodel.compile(loss="binary_crossentropy",\n                    optimizer="r
msprop",\n                    metrics=["accuracy"])\n'

```

Training the regularized convnet

In [19]:

```

"""
callbacks = [
    keras.callbacks.ModelCheckpoint(

```

```

        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=100,
    validation_data=validation_dataset,
    callbacks=callbacks)
"""

```

Out[19]:

```

'\ncallbacks = [\n    keras.callbacks.ModelCheckpoint(\n        filepath="convnet_from_scratch_with_augmentation.keras",\n        save_best_only=True,\n        monitor="val_loss")\n]\nhistory = model.fit(\n    train_dataset,\n    epochs=100,\n    validation_data=validation_dataset,\n    callbacks=callbacks)\n\n'

```

Evaluating the model on the test set

In [20]:

```

"""
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
"""

```

Out[20]:

```

'\ntest_model = keras.models.load_model(\n    "convnet_from_scratch_with_augmentation.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n\n'

```

Leveraging a pretrained model

Feature extraction with a pretrained model

Instantiating the VGG16 convolutional base

In [21]:

```

from tensorflow import keras # import keras
from tensorflow.keras import layers
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False,
    input_shape=(180, 180, 3))

```

In [22]:

```

conv_base.summary()
Model: "vgg16"

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 180, 180, 3)]	0

block1_conv1 (Conv2D)	(None, 180, 180, 64)	1792
block1_conv2 (Conv2D)	(None, 180, 180, 64)	36928
block1_pool (MaxPooling2D)	(None, 90, 90, 64)	0
block2_conv1 (Conv2D)	(None, 90, 90, 128)	73856
block2_conv2 (Conv2D)	(None, 90, 90, 128)	147584
block2_pool (MaxPooling2D)	(None, 45, 45, 128)	0
block3_conv1 (Conv2D)	(None, 45, 45, 256)	295168
block3_conv2 (Conv2D)	(None, 45, 45, 256)	590080
block3_conv3 (Conv2D)	(None, 45, 45, 256)	590080
block3_pool (MaxPooling2D)	(None, 22, 22, 256)	0
block4_conv1 (Conv2D)	(None, 22, 22, 512)	1180160
block4_conv2 (Conv2D)	(None, 22, 22, 512)	2359808
block4_conv3 (Conv2D)	(None, 22, 22, 512)	2359808
block4_pool (MaxPooling2D)	(None, 11, 11, 512)	0
block5_conv1 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv2 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv3 (Conv2D)	(None, 11, 11, 512)	2359808
block5_pool (MaxPooling2D)	(None, 5, 5, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

Fast feature extraction without data augmentation

Extracting the VGG16 features and corresponding labels

In [23]:

```
import numpy as np

def get_features_and_labels(dataset):
    all_features = []
    all_labels = []
    for images, labels in dataset:
```

```

        preprocessed_images =
keras.applications.vgg16.preprocess_input(images)
        features = conv_base.predict(preprocessed_images)
        all_features.append(features)
        all_labels.append(labels)
    return np.concatenate(all_features), np.concatenate(all_labels)

train_features, train_labels = get_features_and_labels(train_dataset)
val_features, val_labels = get_features_and_labels(validation_dataset)
test_features, test_labels = get_features_and_labels(test_dataset)

```

In [24]:

```
train_features.shape
```

Out[24]:

```
(2000, 5, 5, 512)
```

Defining and training the densely connected classifier

In [25]:

```

inputs = keras.Input(shape=(5, 5, 512))
x = layers.Flatten()(inputs)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_features, train_labels,
    epochs=20,
    validation_data=(val_features, val_labels),
    callbacks=callbacks)

Epoch 1/20
63/63 [=====] - 1s 6ms/step - loss: 13.5073 - accuracy: 0.9210 - val_loss: 4.7467 - val_accuracy: 0.9660
Epoch 2/20
63/63 [=====] - 0s 4ms/step - loss: 3.0386 - accuracy: 0.9770 - val_loss: 12.0533 - val_accuracy: 0.9430
Epoch 3/20
63/63 [=====] - 0s 3ms/step - loss: 2.1036 - accuracy: 0.9850 - val_loss: 3.2843 - val_accuracy: 0.9750
Epoch 4/20
63/63 [=====] - 0s 4ms/step - loss: 1.9179 - accuracy: 0.9875 - val_loss: 7.3282 - val_accuracy: 0.9700
Epoch 5/20

```

```

63/63 [=====] - 0s 3ms/step - loss: 0.9677 - accurac
y: 0.9900 - val_loss: 6.0118 - val_accuracy: 0.9650
Epoch 6/20
63/63 [=====] - 0s 3ms/step - loss: 1.2342 - accurac
y: 0.9935 - val_loss: 6.0537 - val_accuracy: 0.9710
Epoch 7/20
63/63 [=====] - 0s 3ms/step - loss: 0.2122 - accurac
y: 0.9975 - val_loss: 6.1084 - val_accuracy: 0.9700
Epoch 8/20
63/63 [=====] - 0s 3ms/step - loss: 0.4881 - accurac
y: 0.9970 - val_loss: 5.7784 - val_accuracy: 0.9710
Epoch 9/20
63/63 [=====] - 0s 3ms/step - loss: 0.0937 - accurac
y: 0.9985 - val_loss: 8.9953 - val_accuracy: 0.9640
Epoch 10/20
63/63 [=====] - 0s 3ms/step - loss: 0.4623 - accurac
y: 0.9970 - val_loss: 5.5047 - val_accuracy: 0.9730
Epoch 11/20
63/63 [=====] - 0s 3ms/step - loss: 0.2748 - accurac
y: 0.9965 - val_loss: 5.4949 - val_accuracy: 0.9760
Epoch 12/20
63/63 [=====] - 0s 3ms/step - loss: 0.4063 - accurac
y: 0.9975 - val_loss: 5.0313 - val_accuracy: 0.9740
Epoch 13/20
63/63 [=====] - 0s 3ms/step - loss: 0.2805 - accurac
y: 0.9965 - val_loss: 6.3703 - val_accuracy: 0.9710
Epoch 14/20
63/63 [=====] - 0s 3ms/step - loss: 0.5061 - accurac
y: 0.9975 - val_loss: 5.6572 - val_accuracy: 0.9760
Epoch 15/20
63/63 [=====] - 0s 3ms/step - loss: 0.0828 - accurac
y: 0.9990 - val_loss: 5.7342 - val_accuracy: 0.9750
Epoch 16/20
63/63 [=====] - 0s 3ms/step - loss: 0.2083 - accurac
y: 0.9975 - val_loss: 5.3191 - val_accuracy: 0.9750
Epoch 17/20
63/63 [=====] - 0s 3ms/step - loss: 0.2813 - accurac
y: 0.9965 - val_loss: 5.0407 - val_accuracy: 0.9760
Epoch 18/20
63/63 [=====] - 0s 3ms/step - loss: 0.0779 - accurac
y: 0.9990 - val_loss: 5.0738 - val_accuracy: 0.9770
Epoch 19/20
63/63 [=====] - 0s 3ms/step - loss: 0.0072 - accurac
y: 0.9990 - val_loss: 5.6584 - val_accuracy: 0.9770
Epoch 20/20
63/63 [=====] - 0s 3ms/step - loss: 3.5128e-17 - acc
uracy: 1.0000 - val_loss: 5.6584 - val_accuracy: 0.9770

```

Plotting the results

In [26]:

```


import matplotlib.pyplot as plt
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]


```

```

loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()

```





Feature extraction together with data augmentation

Instantiating and freezing the VGG16 convolutional base

In [27]:

```
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False)
conv_base.trainable = False
```

Printing the list of trainable weights before and after freezing

In [28]:

```
conv_base.trainable = True
print("This is the number of trainable weights "
      "before freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 26

In [29]:

```
conv_base.trainable = False
print("This is the number of trainable weights "
      "after freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights after freezing the conv base: 0

Adding a data augmentation stage and a classifier to the convolutional base

In [30]:

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
```



```

    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

In [31]:

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction_with_data_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=50,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/50
63/63 [=====] - 5s 56ms/step - loss: 21.2061 - accuracy: 0.8910 - val_loss: 22.9828 - val_accuracy: 0.8940
Epoch 2/50
63/63 [=====] - 3s 52ms/step - loss: 7.0033 - accuracy: 0.9445 - val_loss: 5.7581 - val_accuracy: 0.9620
Epoch 3/50
63/63 [=====] - 3s 53ms/step - loss: 6.5587 - accuracy: 0.9570 - val_loss: 5.6807 - val_accuracy: 0.9590
Epoch 4/50
63/63 [=====] - 4s 56ms/step - loss: 4.4935 - accuracy: 0.9605 - val_loss: 6.0700 - val_accuracy: 0.9640
Epoch 5/50
63/63 [=====] - 4s 56ms/step - loss: 3.2817 - accuracy: 0.9675 - val_loss: 6.1735 - val_accuracy: 0.9650
Epoch 6/50
63/63 [=====] - 3s 52ms/step - loss: 5.4043 - accuracy: 0.9590 - val_loss: 7.8655 - val_accuracy: 0.9640
Epoch 7/50
63/63 [=====] - 3s 53ms/step - loss: 4.8845 - accuracy: 0.9620 - val_loss: 2.8868 - val_accuracy: 0.9760
Epoch 8/50
63/63 [=====] - 3s 53ms/step - loss: 2.2545 - accuracy: 0.9765 - val_loss: 5.1264 - val_accuracy: 0.9700
Epoch 9/50

63/63 [=====] - 3s 52ms/step - loss: 2.3987 - accuracy: 0.9770 - val_loss: 4.8956 - val_accuracy: 0.9740
Epoch 10/50
63/63 [=====] - 4s 57ms/step - loss: 2.7642 - accuracy: 0.9730 - val_loss: 4.4262 - val_accuracy: 0.9770
Epoch 11/50
63/63 [=====] - 3s 52ms/step - loss: 2.7902 - accuracy: 0.9755 - val_loss: 3.1060 - val_accuracy: 0.9770
Epoch 12/50
63/63 [=====] - 3s 53ms/step - loss: 2.5042 - accuracy: 0.9755 - val_loss: 4.9537 - val_accuracy: 0.9680
Epoch 13/50
63/63 [=====] - 3s 52ms/step - loss: 1.4807 - accuracy: 0.9795 - val_loss: 3.4375 - val_accuracy: 0.9790
Epoch 14/50
63/63 [=====] - 3s 52ms/step - loss: 1.5025 - accuracy: 0.9840 - val_loss: 3.9441 - val_accuracy: 0.9740
Epoch 15/50
63/63 [=====] - 3s 52ms/step - loss: 1.9034 - accuracy: 0.9805 - val_loss: 3.6990 - val_accuracy: 0.9730
Epoch 16/50
63/63 [=====] - 3s 52ms/step - loss: 2.0529 - accuracy: 0.9770 - val_loss: 4.0766 - val_accuracy: 0.9780
Epoch 17/50
63/63 [=====] - 3s 52ms/step - loss: 1.4646 - accuracy: 0.9790 - val_loss: 3.3099 - val_accuracy: 0.9770
Epoch 18/50
63/63 [=====] - 3s 52ms/step - loss: 1.4809 - accuracy: 0.9820 - val_loss: 3.2118 - val_accuracy: 0.9720
Epoch 19/50
63/63 [=====] - 3s 52ms/step - loss: 1.1507 - accuracy: 0.9815 - val_loss: 2.6941 - val_accuracy: 0.9750
Epoch 20/50
63/63 [=====] - 3s 53ms/step - loss: 1.6299 - accuracy: 0.9805 - val_loss: 2.7439 - val_accuracy: 0.9760
Epoch 21/50
63/63 [=====] - 3s 53ms/step - loss: 1.1083 - accuracy: 0.9845 - val_loss: 3.9618 - val_accuracy: 0.9790
Epoch 22/50
63/63 [=====] - 3s 52ms/step - loss: 1.3424 - accuracy: 0.9835 - val_loss: 3.0976 - val_accuracy: 0.9810
Epoch 23/50
63/63 [=====] - 3s 52ms/step - loss: 1.3167 - accuracy: 0.9795 - val_loss: 3.8232 - val_accuracy: 0.9720
Epoch 24/50
63/63 [=====] - 3s 53ms/step - loss: 1.3782 - accuracy: 0.9810 - val_loss: 2.3537 - val_accuracy: 0.9780
Epoch 25/50
63/63 [=====] - 3s 53ms/step - loss: 0.9867 - accuracy: 0.9845 - val_loss: 2.9392 - val_accuracy: 0.9770
Epoch 26/50
63/63 [=====] - 3s 53ms/step - loss: 0.9766 - accuracy: 0.9825 - val_loss: 2.4336 - val_accuracy: 0.9780

Epoch 27/50
63/63 [=====] - 3s 52ms/step - loss: 1.0308 - accuracy: 0.9830 - val_loss: 3.6002 - val_accuracy: 0.9720
Epoch 28/50
63/63 [=====] - 3s 52ms/step - loss: 0.6114 - accuracy: 0.9850 - val_loss: 2.4190 - val_accuracy: 0.9810
Epoch 29/50
63/63 [=====] - 3s 52ms/step - loss: 0.7387 - accuracy: 0.9855 - val_loss: 4.1295 - val_accuracy: 0.9660
Epoch 30/50
63/63 [=====] - 3s 52ms/step - loss: 0.5670 - accuracy: 0.9845 - val_loss: 3.4808 - val_accuracy: 0.9740
Epoch 31/50
63/63 [=====] - 3s 53ms/step - loss: 0.7438 - accuracy: 0.9855 - val_loss: 2.7265 - val_accuracy: 0.9750
Epoch 32/50
63/63 [=====] - 3s 52ms/step - loss: 0.6358 - accuracy: 0.9880 - val_loss: 4.9777 - val_accuracy: 0.9630
Epoch 33/50
63/63 [=====] - 3s 52ms/step - loss: 1.0329 - accuracy: 0.9845 - val_loss: 2.0601 - val_accuracy: 0.9770
Epoch 34/50
63/63 [=====] - 3s 52ms/step - loss: 0.5896 - accuracy: 0.9895 - val_loss: 2.1925 - val_accuracy: 0.9800
Epoch 35/50
63/63 [=====] - 3s 52ms/step - loss: 0.9808 - accuracy: 0.9855 - val_loss: 2.6108 - val_accuracy: 0.9790
Epoch 36/50
63/63 [=====] - 4s 53ms/step - loss: 0.8450 - accuracy: 0.9850 - val_loss: 2.7040 - val_accuracy: 0.9710
Epoch 37/50
63/63 [=====] - 3s 53ms/step - loss: 0.5562 - accuracy: 0.9875 - val_loss: 2.3025 - val_accuracy: 0.9730
Epoch 38/50
63/63 [=====] - 3s 51ms/step - loss: 1.2197 - accuracy: 0.9820 - val_loss: 2.5752 - val_accuracy: 0.9760
Epoch 39/50
63/63 [=====] - 3s 52ms/step - loss: 0.6530 - accuracy: 0.9890 - val_loss: 3.0409 - val_accuracy: 0.9690
Epoch 40/50
63/63 [=====] - 3s 52ms/step - loss: 0.3301 - accuracy: 0.9940 - val_loss: 2.8597 - val_accuracy: 0.9800
Epoch 41/50
63/63 [=====] - 3s 52ms/step - loss: 0.5592 - accuracy: 0.9895 - val_loss: 2.3041 - val_accuracy: 0.9750
Epoch 42/50
63/63 [=====] - 3s 52ms/step - loss: 0.5874 - accuracy: 0.9920 - val_loss: 2.4542 - val_accuracy: 0.9760
Epoch 43/50
63/63 [=====] - 3s 52ms/step - loss: 0.3692 - accuracy: 0.9905 - val_loss: 2.3337 - val_accuracy: 0.9780
Epoch 44/50

```

63/63 [=====] - 3s 53ms/step - loss: 0.9109 - accuracy: 0.9860 - val_loss: 2.7692 - val_accuracy: 0.9780
Epoch 45/50
63/63 [=====] - 3s 52ms/step - loss: 0.8140 - accuracy: 0.9860 - val_loss: 2.4054 - val_accuracy: 0.9750
Epoch 46/50
63/63 [=====] - 3s 53ms/step - loss: 0.7053 - accuracy: 0.9885 - val_loss: 2.5727 - val_accuracy: 0.9790
Epoch 47/50
63/63 [=====] - 3s 53ms/step - loss: 0.6289 - accuracy: 0.9865 - val_loss: 2.2830 - val_accuracy: 0.9760
Epoch 48/50
63/63 [=====] - 3s 52ms/step - loss: 0.5192 - accuracy: 0.9885 - val_loss: 2.3128 - val_accuracy: 0.9780
Epoch 49/50
63/63 [=====] - 3s 52ms/step - loss: 0.5046 - accuracy: 0.9905 - val_loss: 2.6822 - val_accuracy: 0.9770
Epoch 50/50
63/63 [=====] - 3s 53ms/step - loss: 0.8204 - accuracy: 0.9860 - val_loss: 2.0805 - val_accuracy: 0.9790

```

Evaluating the model on the test set

In [32]:

```

test_model = keras.models.load_model(
    "feature_extraction_with_data_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 32ms/step - loss: 2.4537 - accuracy: 0.9770
Test accuracy: 0.977

```

Fine-tuning a pretrained model

In [33]:

```

conv_base.summary()
Model: "vgg16"

```

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584

block2_pool	(MaxPooling2D)	(None, None, None, 128)	0
block3_conv1	(Conv2D)	(None, None, None, 256)	295168
block3_conv2	(Conv2D)	(None, None, None, 256)	590080
block3_conv3	(Conv2D)	(None, None, None, 256)	590080
block3_pool	(MaxPooling2D)	(None, None, None, 256)	0
block4_conv1	(Conv2D)	(None, None, None, 512)	1180160
block4_conv2	(Conv2D)	(None, None, None, 512)	2359808
block4_conv3	(Conv2D)	(None, None, None, 512)	2359808
block4_pool	(MaxPooling2D)	(None, None, None, 512)	0
block5_conv1	(Conv2D)	(None, None, None, 512)	2359808
block5_conv2	(Conv2D)	(None, None, None, 512)	2359808
block5_conv3	(Conv2D)	(None, None, None, 512)	2359808
block5_pool	(MaxPooling2D)	(None, None, None, 512)	0

=====

Total params: 14,714,688
Trainable params: 0
Non-trainable params: 14,714,688

Freezing all layers until the fourth from the last

In [34]:

```
conv_base.trainable = True
for layer in conv_base.layers[:-4]:
    layer.trainable = False
```

Fine-tuning the model

In [35]:

```
model.compile(loss="binary_crossentropy",
              optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
```

callbacks=callbacks)

Epoch 1/30

63/63 [=====] - 5s 61ms/step - loss: 0.3013 - accuracy: 0.9920 - val_loss: 2.3283 - val_accuracy: 0.9720

Epoch 2/30

63/63 [=====] - 4s 56ms/step - loss: 0.4117 - accuracy: 0.9850 - val_loss: 2.0034 - val_accuracy: 0.9790

Epoch 3/30

63/63 [=====] - 4s 56ms/step - loss: 0.4557 - accuracy: 0.9890 - val_loss: 2.1317 - val_accuracy: 0.9780

Epoch 4/30

63/63 [=====] - 4s 56ms/step - loss: 0.3864 - accuracy: 0.9875 - val_loss: 2.2103 - val_accuracy: 0.9800

Epoch 5/30

63/63 [=====] - 4s 57ms/step - loss: 0.2877 - accuracy: 0.9920 - val_loss: 2.3774 - val_accuracy: 0.9760

Epoch 6/30

63/63 [=====] - 4s 57ms/step - loss: 0.2910 - accuracy: 0.9890 - val_loss: 2.4382 - val_accuracy: 0.9740

Epoch 7/30

63/63 [=====] - 4s 56ms/step - loss: 0.1776 - accuracy: 0.9930 - val_loss: 2.0192 - val_accuracy: 0.9780

Epoch 8/30

63/63 [=====] - 4s 56ms/step - loss: 0.3179 - accuracy: 0.9915 - val_loss: 2.1185 - val_accuracy: 0.9750

Epoch 9/30

63/63 [=====] - 4s 56ms/step - loss: 0.2422 - accuracy: 0.9930 - val_loss: 1.8964 - val_accuracy: 0.9770

Epoch 10/30

63/63 [=====] - 4s 56ms/step - loss: 0.2623 - accuracy: 0.9940 - val_loss: 1.5055 - val_accuracy: 0.9820

Epoch 11/30

63/63 [=====] - 4s 56ms/step - loss: 0.2996 - accuracy: 0.9920 - val_loss: 1.2986 - val_accuracy: 0.9820

Epoch 12/30

63/63 [=====] - 4s 56ms/step - loss: 0.0817 - accuracy: 0.9975 - val_loss: 1.7039 - val_accuracy: 0.9800

Epoch 13/30

63/63 [=====] - 4s 56ms/step - loss: 0.2091 - accuracy: 0.9940 - val_loss: 1.7767 - val_accuracy: 0.9820

Epoch 14/30

63/63 [=====] - 4s 56ms/step - loss: 0.2408 - accuracy: 0.9940 - val_loss: 1.6753 - val_accuracy: 0.9810

Epoch 15/30

63/63 [=====] - 4s 56ms/step - loss: 0.2389 - accuracy: 0.9950 - val_loss: 1.4084 - val_accuracy: 0.9800

Epoch 16/30

63/63 [=====] - 4s 56ms/step - loss: 0.1064 - accuracy: 0.9955 - val_loss: 1.6390 - val_accuracy: 0.9790

Epoch 17/30

63/63 [=====] - 4s 56ms/step - loss: 0.0449 - accuracy: 0.9990 - val_loss: 1.6157 - val_accuracy: 0.9790

```

Epoch 18/30
63/63 [=====] - 4s 56ms/step - loss: 0.1464 - accuracy: 0.9930 - val_loss: 1.7780 - val_accuracy: 0.9770
Epoch 19/30
63/63 [=====] - 4s 56ms/step - loss: 0.2795 - accuracy: 0.9955 - val_loss: 1.5480 - val_accuracy: 0.9780
Epoch 20/30
63/63 [=====] - 4s 56ms/step - loss: 0.0916 - accuracy: 0.9970 - val_loss: 1.7212 - val_accuracy: 0.9760
Epoch 21/30
63/63 [=====] - 4s 57ms/step - loss: 0.0743 - accuracy: 0.9965 - val_loss: 1.1029 - val_accuracy: 0.9800
Epoch 22/30
63/63 [=====] - 4s 56ms/step - loss: 0.1037 - accuracy: 0.9960 - val_loss: 1.8676 - val_accuracy: 0.9760
Epoch 23/30
63/63 [=====] - 4s 57ms/step - loss: 0.1488 - accuracy: 0.9940 - val_loss: 1.5035 - val_accuracy: 0.9810
Epoch 24/30
63/63 [=====] - 4s 56ms/step - loss: 0.0420 - accuracy: 0.9970 - val_loss: 1.4502 - val_accuracy: 0.9810
Epoch 25/30
63/63 [=====] - 4s 56ms/step - loss: 0.0781 - accuracy: 0.9965 - val_loss: 1.5331 - val_accuracy: 0.9810
Epoch 26/30
63/63 [=====] - 4s 57ms/step - loss: 0.1049 - accuracy: 0.9955 - val_loss: 1.3614 - val_accuracy: 0.9800
Epoch 27/30
63/63 [=====] - 4s 56ms/step - loss: 0.0205 - accuracy: 0.9985 - val_loss: 1.2473 - val_accuracy: 0.9820
Epoch 28/30
63/63 [=====] - 4s 56ms/step - loss: 0.1799 - accuracy: 0.9940 - val_loss: 1.7278 - val_accuracy: 0.9790
Epoch 29/30
63/63 [=====] - 4s 56ms/step - loss: 0.0545 - accuracy: 0.9980 - val_loss: 1.5502 - val_accuracy: 0.9820
Epoch 30/30
63/63 [=====] - 4s 57ms/step - loss: 0.0587 - accuracy: 0.9950 - val_loss: 2.0134 - val_accuracy: 0.9760

```

In [36]:

```

model = keras.models.load_model("fine_tuning.keras")
test_loss, test_acc = model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 33ms/step - loss: 1.5191 - accuracy: 0.9810
Test accuracy: 0.981

```

Summary

Using a pretrained network with training sample of 3000, a validation sample of 500, and a test sample of 500

Downloading the data

```
#!/unzip -qq '/fs/ess/PGS0333/BA_64061_KSU/data/dogs-vs-cats.zip'
```

In [1]:

```
#!/unzip -qq train.zip
```

In [2]:

Copying images to training, validation, and test directories

```
import os, shutil, pathlib

original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index,
end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=2999)
make_subset("validation", start_index=3000, end_index=3499)
make_subset("test", start_index=3500, end_index=3999)
```

In [3]:

Building the model

Instantiating a small convnet for dogs vs. cats classification

```
"""
from tensorflow import keras
from tensorflow.keras import layers
inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
```

In [4]:


```

x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
"""

```

Out[4]:

```

'\nfrom tensorflow import keras\nfrom tensorflow.keras import layers\ninputs
= keras.Input(shape=(180, 180, 3))\nx = layers.Rescaling(1./255)(inputs)\nx =
layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)\nx = layers.M
axPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filters=64, kernel_size=3, act
ivation="relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2
D(filters=128, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling2D
(pool_size=2)(x)\nx = layers.Conv2D(filters=256, kernel_size=3, activation="r
elu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filters=
256, kernel_size=3, activation="relu")(x)\nx = layers.Flatten()(x)\noutputs =
layers.Dense(1, activation="sigmoid")(x)\nmodel = keras.Model(inputs=inputs,
outputs=outputs)\n'

```

In [5]:

```

# model.summary()

```

Configuring the model for training

In [6]:

```

"""
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
"""

```

Out[6]:

```

'\nmodel.compile(loss="binary_crossentropy",\n
optimizer="rmspro
p",\n
metrics=["accuracy"])\n'

```

Data preprocessing

Using image_dataset_from_directory to read images

In [7]:

```

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",

```

```
image_size=(180, 180),  
batch_size=32)
```

```
Found 5998 files belonging to 2 classes.  
Found 998 files belonging to 2 classes.  
Found 998 files belonging to 2 classes.
```

In [8]:

```
"""  
import numpy as np  
import tensorflow as tf  
random_numbers = np.random.normal(size=(1000, 16))  
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)  
"""
```

Out[8]:

```
'\nimport numpy as np\nimport tensorflow as tf\nrandom_numbers = np.random.normal(size=(1000, 16))\ndataset = tf.data.Dataset.from_tensor_slices(random_numbers)\n'
```

In [9]:

```
"""  
for i, element in enumerate(dataset):  
    print(element.shape)  
    if i >= 2:  
        break  
"""
```

Out[9]:

```
'\nfor i, element in enumerate(dataset):\n    print(element.shape)\n    if i >= 2:\n        break\n'
```

In [10]:

```
"""  
batched_dataset = dataset.batch(32)  
for i, element in enumerate(batched_dataset):  
    print(element.shape)  
    if i >= 2:  
        break  
"""
```

Out[10]:

```
'\nbatched_dataset = dataset.batch(32)\nfor i, element in enumerate(batched_dataset):\n    print(element.shape)\n    if i >= 2:\n        break\n'
```

In [11]:

```
"""  
reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))  
for i, element in enumerate(reshaped_dataset):  
    print(element.shape)  
    if i >= 2:  
        break  
"""
```

Out[11]:

```
'\nreshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))\nfor i, element in enumerate(reshaped_dataset):\n    print(element.shape)\n    if i >= 2:\n        break\n'
```

Displaying the shapes of the data and labels yielded by the Dataset

In [12]:

```
"""
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    break
"""
```

Out[12]:

```
'\nfor data_batch, labels_batch in train_dataset:\n    print("data batch shape:", data_batch.shape)\n    print("labels batch shape:", labels_batch.shape)\n    break\n'
```

Fitting the model using a Dataset

In [13]:

```
"""
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)
"""
```

Out[13]:

```
'\ncallbacks = [\n    keras.callbacks.ModelCheckpoint(\n        filepath="convnet_from_scratch.keras",\n        save_best_only=True,\n        monitor="val_loss")\n]\nhistory = model.fit(\n    train_dataset,\n    epochs=30,\n    validation_data=validation_dataset,\n    callbacks=callbacks)\n'
```

Displaying curves of loss and accuracy during training

In [14]:

```
"""
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
"""
```

```
plt.show()
"""
```

Out[14]:

```
'\nimport matplotlib.pyplot as plt\naccuracy = history.history["accuracy"]\nval_accuracy = history.history["val_accuracy"]\nloss = history.history["loss"]\nval_loss = history.history["val_loss"]\nepochs = range(1, len(accuracy) + 1)\nplt.plot(epochs, accuracy, "bo", label="Training accuracy")\nplt.plot(epochs, val_accuracy, "b", label="Validation accuracy")\nplt.title("Training and validation accuracy")\nplt.legend()\nplt.figure()\nplt.plot(epochs, loss, "bo", label="Training loss")\nplt.plot(epochs, val_loss, "b", label="Validation loss")\nplt.title("Training and validation loss")\nplt.legend()\nplt.show()\n'
```

Evaluating the model on the test set

In [15]:

```
"""
test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
"""
```

Out[15]:

```
'\ntest_model = keras.models.load_model("convnet_from_scratch.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n'
```

Using data augmentation

Define a data augmentation stage to add to an image model

In [16]:

```
"""
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
"""
```

Out[16]:

```
'\ndata_augmentation = keras.Sequential(\n    [\n        layers.RandomFlip("horizontal"),\n        layers.RandomRotation(0.1),\n        layers.RandomZoom(0.2),\n    ]\n)\n'
```

Displaying some randomly augmented training images

In [17]:

```
"""
plt.figure(figsize=(10, 10))
for images, _ in train_dataset.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
```

```

plt.imshow(augmented_images[0].numpy().astype("uint8"))
plt.axis("off")
"""

```

Out[17]:

```

'\nplt.figure(figsize=(10, 10))\nfor images, _ in train_dataset.take(1):\n
  for i in range(9):\n      augmented_images = data_augmentation(images)\n
      ax = plt.subplot(3, 3, i + 1)\n      plt.imshow(augmented_images[0].n
umpy().astype("uint8"))\n      plt.axis("off")\n'

```

Defining a new convnet that includes image augmentation and dropout

In [18]:

```

"""
inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
"""

```

Out[18]:

```

'\ninputs = keras.Input(shape=(180, 180, 3))\nx = data_augmentation(inputs)\n
x = layers.Rescaling(1./255)(x)\nx = layers.Conv2D(filters=32, kernel_size=3,
  activation="relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.C
onv2D(filters=64, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling
2D(pool_size=2)(x)\nx = layers.Conv2D(filters=128, kernel_size=3, activation=
"relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filter
s=256, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling2D(pool_siz
e=2)(x)\nx = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)\
nx = layers.Flatten()(x)\nx = layers.Dropout(0.5)(x)\noutputs = layers.Dense
(1, activation="sigmoid")(x)\nmodel = keras.Model(inputs=inputs, outputs=outp
uts)\n\nmodel.compile(loss="binary_crossentropy",\n                    optimizer="r
msprop",\n                    metrics=["accuracy"])\n'

```

Training the regularized convnet

In [19]:

```

"""
callbacks = [
    keras.callbacks.ModelCheckpoint(

```

```

        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=100,
    validation_data=validation_dataset,
    callbacks=callbacks)
"""

```

Out[19]:

```

'\ncallbacks = [\n    keras.callbacks.ModelCheckpoint(\n        filepath="convnet_from_scratch_with_augmentation.keras",\n        save_best_only=True,\n        monitor="val_loss")\n]\nhistory = model.fit(\n    train_dataset,\n    epochs=100,\n    validation_data=validation_dataset,\n    callbacks=callbacks)\n\n'

```

Evaluating the model on the test set

In [20]:

```

"""
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
"""

```

Out[20]:

```

'\ntest_model = keras.models.load_model(\n    "convnet_from_scratch_with_augmentation.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n\n'

```

Leveraging a pretrained model

Feature extraction with a pretrained model

Instantiating the VGG16 convolutional base

In [21]:

```

from tensorflow import keras # import keras
from tensorflow.keras import layers
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False,
    input_shape=(180, 180, 3))

```

In [22]:

```

conv_base.summary()
Model: "vgg16"

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 180, 180, 3)]	0

block1_conv1 (Conv2D)	(None, 180, 180, 64)	1792
block1_conv2 (Conv2D)	(None, 180, 180, 64)	36928
block1_pool (MaxPooling2D)	(None, 90, 90, 64)	0
block2_conv1 (Conv2D)	(None, 90, 90, 128)	73856
block2_conv2 (Conv2D)	(None, 90, 90, 128)	147584
block2_pool (MaxPooling2D)	(None, 45, 45, 128)	0
block3_conv1 (Conv2D)	(None, 45, 45, 256)	295168
block3_conv2 (Conv2D)	(None, 45, 45, 256)	590080
block3_conv3 (Conv2D)	(None, 45, 45, 256)	590080
block3_pool (MaxPooling2D)	(None, 22, 22, 256)	0
block4_conv1 (Conv2D)	(None, 22, 22, 512)	1180160
block4_conv2 (Conv2D)	(None, 22, 22, 512)	2359808
block4_conv3 (Conv2D)	(None, 22, 22, 512)	2359808
block4_pool (MaxPooling2D)	(None, 11, 11, 512)	0
block5_conv1 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv2 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv3 (Conv2D)	(None, 11, 11, 512)	2359808
block5_pool (MaxPooling2D)	(None, 5, 5, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

Fast feature extraction without data augmentation

Extracting the VGG16 features and corresponding labels

In [23]:

```
import numpy as np

def get_features_and_labels(dataset):
    all_features = []
    all_labels = []
    for images, labels in dataset:
```

```

        preprocessed_images =
keras.applications.vgg16.preprocess_input(images)
        features = conv_base.predict(preprocessed_images)
        all_features.append(features)
        all_labels.append(labels)
    return np.concatenate(all_features), np.concatenate(all_labels)

train_features, train_labels = get_features_and_labels(train_dataset)
val_features, val_labels = get_features_and_labels(validation_dataset)
test_features, test_labels = get_features_and_labels(test_dataset)

```

In [24]:

```
train_features.shape
```

Out[24]:

```
(5998, 5, 5, 512)
```

Defining and training the densely connected classifier

In [25]:

```

inputs = keras.Input(shape=(5, 5, 512))
x = layers.Flatten()(inputs)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_features, train_labels,
    epochs=20,
    validation_data=(val_features, val_labels),
    callbacks=callbacks)

Epoch 1/20
188/188 [=====] - 1s 4ms/step - loss: 8.3404 - accur
acy: 0.9505 - val_loss: 5.5009 - val_accuracy: 0.9739
Epoch 2/20
188/188 [=====] - 1s 3ms/step - loss: 3.3456 - accur
acy: 0.9800 - val_loss: 5.1963 - val_accuracy: 0.9800
Epoch 3/20
188/188 [=====] - 0s 3ms/step - loss: 2.2134 - accur
acy: 0.9860 - val_loss: 4.7894 - val_accuracy: 0.9760
Epoch 4/20
188/188 [=====] - 0s 2ms/step - loss: 1.6404 - accur
acy: 0.9892 - val_loss: 3.9754 - val_accuracy: 0.9810
Epoch 5/20

```



```

188/188 [=====] - 1s 3ms/step - loss: 1.4981 - accur
acy: 0.9888 - val_loss: 5.3861 - val_accuracy: 0.9719
Epoch 6/20
188/188 [=====] - 0s 2ms/step - loss: 0.8559 - accur
acy: 0.9938 - val_loss: 3.9264 - val_accuracy: 0.9820
Epoch 7/20
188/188 [=====] - 0s 2ms/step - loss: 0.6068 - accur
acy: 0.9955 - val_loss: 6.1616 - val_accuracy: 0.9780
Epoch 8/20
188/188 [=====] - 0s 2ms/step - loss: 0.2739 - accur
acy: 0.9965 - val_loss: 4.5638 - val_accuracy: 0.9760
Epoch 9/20
188/188 [=====] - 0s 2ms/step - loss: 0.2804 - accur
acy: 0.9970 - val_loss: 4.9678 - val_accuracy: 0.9790
Epoch 10/20
188/188 [=====] - 0s 2ms/step - loss: 0.2552 - accur
acy: 0.9973 - val_loss: 7.0815 - val_accuracy: 0.9749
Epoch 11/20
188/188 [=====] - 0s 2ms/step - loss: 0.2234 - accur
acy: 0.9972 - val_loss: 5.6555 - val_accuracy: 0.9749
Epoch 12/20
188/188 [=====] - 0s 2ms/step - loss: 0.2629 - accur
acy: 0.9972 - val_loss: 6.1986 - val_accuracy: 0.9760
Epoch 13/20
188/188 [=====] - 0s 2ms/step - loss: 0.1071 - accur
acy: 0.9987 - val_loss: 5.3456 - val_accuracy: 0.9739
Epoch 14/20
188/188 [=====] - 0s 2ms/step - loss: 0.2354 - accur
acy: 0.9978 - val_loss: 4.6037 - val_accuracy: 0.9760
Epoch 15/20
188/188 [=====] - 0s 2ms/step - loss: 0.3182 - accur
acy: 0.9980 - val_loss: 5.5216 - val_accuracy: 0.9760
Epoch 16/20
188/188 [=====] - 0s 2ms/step - loss: 0.1429 - accur
acy: 0.9977 - val_loss: 5.7715 - val_accuracy: 0.9689
Epoch 17/20
188/188 [=====] - 0s 2ms/step - loss: 0.0636 - accur
acy: 0.9985 - val_loss: 5.8755 - val_accuracy: 0.9749
Epoch 18/20
188/188 [=====] - 0s 2ms/step - loss: 0.1503 - accur
acy: 0.9985 - val_loss: 6.7080 - val_accuracy: 0.9770
Epoch 19/20
188/188 [=====] - 0s 2ms/step - loss: 0.0935 - accur
acy: 0.9987 - val_loss: 5.6859 - val_accuracy: 0.9749
Epoch 20/20
188/188 [=====] - 0s 2ms/step - loss: 0.0171 - accur
acy: 0.9993 - val_loss: 5.6052 - val_accuracy: 0.9780

```

Plotting the results

In [26]:

```


import matplotlib.pyplot as plt
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]


```

```

loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()

```





Feature extraction together with data augmentation

Instantiating and freezing the VGG16 convolutional base

In [27]:

```
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False)
conv_base.trainable = False
```

Printing the list of trainable weights before and after freezing

In [28]:

```
conv_base.trainable = True
print("This is the number of trainable weights "
      "before freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 26

In [29]:

```
conv_base.trainable = False
print("This is the number of trainable weights "
      "after freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights after freezing the conv base: 0

Adding a data augmentation stage and a classifier to the convolutional base

In [30]:

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
```

```

    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

In [31]:

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction_with_data_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=50,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

```

Epoch 1/50
188/188 [=====] - 9s 42ms/step - loss: 11.1415 - accu
racy: 0.9263 - val_loss: 2.8871 - val_accuracy: 0.9739
Epoch 2/50
188/188 [=====] - 8s 39ms/step - loss: 6.5339 - accu
racy: 0.9498 - val_loss: 4.5935 - val_accuracy: 0.9669
Epoch 3/50
188/188 [=====] - 8s 40ms/step - loss: 4.3118 - accu
racy: 0.9548 - val_loss: 1.8875 - val_accuracy: 0.9780
Epoch 4/50
188/188 [=====] - 8s 40ms/step - loss: 3.2748 - accu
racy: 0.9592 - val_loss: 3.2645 - val_accuracy: 0.9709
Epoch 5/50
188/188 [=====] - 8s 40ms/step - loss: 2.1827 - accu
racy: 0.9608 - val_loss: 1.2957 - val_accuracy: 0.9780
Epoch 6/50
188/188 [=====] - 8s 40ms/step - loss: 1.7407 - accu
racy: 0.9615 - val_loss: 0.8285 - val_accuracy: 0.9810
Epoch 7/50
188/188 [=====] - 8s 40ms/step - loss: 1.1840 - accu
racy: 0.9627 - val_loss: 0.6427 - val_accuracy: 0.9820
Epoch 8/50
188/188 [=====] - 8s 40ms/step - loss: 0.8556 - accu
racy: 0.9668 - val_loss: 0.3899 - val_accuracy: 0.9830
Epoch 9/50

```

188/188 [=====] - 8s 40ms/step - loss: 0.7557 - accuracy: 0.9648 - val_loss: 0.6245 - val_accuracy: 0.9800
Epoch 10/50
188/188 [=====] - 8s 39ms/step - loss: 0.6924 - accuracy: 0.9673 - val_loss: 0.6261 - val_accuracy: 0.9810
Epoch 11/50
188/188 [=====] - 8s 39ms/step - loss: 0.6886 - accuracy: 0.9693 - val_loss: 0.5560 - val_accuracy: 0.9830
Epoch 12/50
188/188 [=====] - 8s 39ms/step - loss: 0.6191 - accuracy: 0.9708 - val_loss: 0.8211 - val_accuracy: 0.9770
Epoch 13/50
188/188 [=====] - 8s 39ms/step - loss: 0.6507 - accuracy: 0.9732 - val_loss: 1.7536 - val_accuracy: 0.9559
Epoch 14/50
188/188 [=====] - 8s 40ms/step - loss: 0.6279 - accuracy: 0.9713 - val_loss: 0.8081 - val_accuracy: 0.9729
Epoch 15/50
188/188 [=====] - 8s 40ms/step - loss: 0.6727 - accuracy: 0.9737 - val_loss: 0.5705 - val_accuracy: 0.9820
Epoch 16/50
188/188 [=====] - 8s 39ms/step - loss: 0.6008 - accuracy: 0.9743 - val_loss: 0.7119 - val_accuracy: 0.9850
Epoch 17/50
188/188 [=====] - 8s 40ms/step - loss: 0.6891 - accuracy: 0.9707 - val_loss: 0.9312 - val_accuracy: 0.9800
Epoch 18/50
188/188 [=====] - 8s 39ms/step - loss: 0.6361 - accuracy: 0.9748 - val_loss: 0.7632 - val_accuracy: 0.9770
Epoch 19/50
188/188 [=====] - 8s 40ms/step - loss: 0.6982 - accuracy: 0.9713 - val_loss: 0.7708 - val_accuracy: 0.9810
Epoch 20/50
188/188 [=====] - 8s 40ms/step - loss: 0.5491 - accuracy: 0.9773 - val_loss: 0.6233 - val_accuracy: 0.9830
Epoch 21/50
188/188 [=====] - 8s 39ms/step - loss: 0.6702 - accuracy: 0.9740 - val_loss: 1.2653 - val_accuracy: 0.9749
Epoch 22/50
188/188 [=====] - 8s 40ms/step - loss: 0.6501 - accuracy: 0.9743 - val_loss: 0.6579 - val_accuracy: 0.9860
Epoch 23/50
188/188 [=====] - 8s 40ms/step - loss: 0.7387 - accuracy: 0.9732 - val_loss: 0.8399 - val_accuracy: 0.9840
Epoch 24/50
188/188 [=====] - 8s 40ms/step - loss: 0.7319 - accuracy: 0.9733 - val_loss: 0.6483 - val_accuracy: 0.9850
Epoch 25/50
188/188 [=====] - 8s 39ms/step - loss: 0.5826 - accuracy: 0.9788 - val_loss: 1.0020 - val_accuracy: 0.9810
Epoch 26/50
188/188 [=====] - 8s 41ms/step - loss: 0.5837 - accuracy: 0.9802 - val_loss: 0.7137 - val_accuracy: 0.9840

Epoch 27/50
188/188 [=====] - 8s 41ms/step - loss: 0.6349 - accuracy: 0.9785 - val_loss: 0.6639 - val_accuracy: 0.9810
Epoch 28/50
188/188 [=====] - 8s 41ms/step - loss: 0.7112 - accuracy: 0.9770 - val_loss: 0.8520 - val_accuracy: 0.9860
Epoch 29/50
188/188 [=====] - 8s 40ms/step - loss: 0.5568 - accuracy: 0.9805 - val_loss: 1.0098 - val_accuracy: 0.9850
Epoch 30/50
188/188 [=====] - 8s 40ms/step - loss: 0.6739 - accuracy: 0.9785 - val_loss: 1.0824 - val_accuracy: 0.9760
Epoch 31/50
188/188 [=====] - 8s 40ms/step - loss: 0.6461 - accuracy: 0.9795 - val_loss: 0.8808 - val_accuracy: 0.9850
Epoch 32/50
188/188 [=====] - 8s 40ms/step - loss: 0.8002 - accuracy: 0.9750 - val_loss: 0.9471 - val_accuracy: 0.9820
Epoch 33/50
188/188 [=====] - 8s 40ms/step - loss: 0.7060 - accuracy: 0.9770 - val_loss: 0.7227 - val_accuracy: 0.9840
Epoch 34/50
188/188 [=====] - 8s 40ms/step - loss: 0.7301 - accuracy: 0.9782 - val_loss: 0.8728 - val_accuracy: 0.9820
Epoch 35/50
188/188 [=====] - 8s 40ms/step - loss: 0.5756 - accuracy: 0.9798 - val_loss: 0.8718 - val_accuracy: 0.9820
Epoch 36/50
188/188 [=====] - 8s 40ms/step - loss: 0.6341 - accuracy: 0.9790 - val_loss: 0.8387 - val_accuracy: 0.9830
Epoch 37/50
188/188 [=====] - 8s 40ms/step - loss: 0.7626 - accuracy: 0.9782 - val_loss: 0.8807 - val_accuracy: 0.9830
Epoch 38/50
188/188 [=====] - 8s 40ms/step - loss: 0.7786 - accuracy: 0.9767 - val_loss: 0.7706 - val_accuracy: 0.9840
Epoch 39/50
188/188 [=====] - 8s 40ms/step - loss: 0.7105 - accuracy: 0.9780 - val_loss: 0.9285 - val_accuracy: 0.9820
Epoch 40/50
188/188 [=====] - 8s 40ms/step - loss: 0.7216 - accuracy: 0.9783 - val_loss: 0.9317 - val_accuracy: 0.9830
Epoch 41/50
188/188 [=====] - 8s 40ms/step - loss: 0.7098 - accuracy: 0.9780 - val_loss: 0.9825 - val_accuracy: 0.9810
Epoch 42/50
188/188 [=====] - 8s 40ms/step - loss: 0.6639 - accuracy: 0.9813 - val_loss: 0.9120 - val_accuracy: 0.9850
Epoch 43/50
188/188 [=====] - 8s 40ms/step - loss: 0.6375 - accuracy: 0.9815 - val_loss: 1.2511 - val_accuracy: 0.9840
Epoch 44/50

```

188/188 [=====] - 8s 40ms/step - loss: 0.7082 - accu
racy: 0.9785 - val_loss: 0.9283 - val_accuracy: 0.9870
Epoch 45/50
188/188 [=====] - 8s 40ms/step - loss: 0.6106 - accu
racy: 0.9845 - val_loss: 0.8911 - val_accuracy: 0.9850
Epoch 46/50
188/188 [=====] - 8s 40ms/step - loss: 0.6737 - accu
racy: 0.9808 - val_loss: 1.0990 - val_accuracy: 0.9810
Epoch 47/50
188/188 [=====] - 8s 40ms/step - loss: 0.7285 - accu
racy: 0.9787 - val_loss: 1.3349 - val_accuracy: 0.9830
Epoch 48/50
188/188 [=====] - 8s 40ms/step - loss: 0.6577 - accu
racy: 0.9830 - val_loss: 1.3496 - val_accuracy: 0.9830
Epoch 49/50
188/188 [=====] - 8s 40ms/step - loss: 0.9656 - accu
racy: 0.9787 - val_loss: 1.0118 - val_accuracy: 0.9800
Epoch 50/50
188/188 [=====] - 8s 40ms/step - loss: 0.7398 - accu
racy: 0.9823 - val_loss: 1.8748 - val_accuracy: 0.9780

```

Evaluating the model on the test set

In [32]:

```

test_model = keras.models.load_model(
    "feature_extraction_with_data_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 33ms/step - loss: 0.7367 - accura
cy: 0.9790
Test accuracy: 0.979

```

Fine-tuning a pretrained model

In [33]:

```

conv_base.summary()
Model: "vgg16"

```

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584

block2_pool	(MaxPooling2D)	(None, None, None, 128)	0
block3_conv1	(Conv2D)	(None, None, None, 256)	295168
block3_conv2	(Conv2D)	(None, None, None, 256)	590080
block3_conv3	(Conv2D)	(None, None, None, 256)	590080
block3_pool	(MaxPooling2D)	(None, None, None, 256)	0
block4_conv1	(Conv2D)	(None, None, None, 512)	1180160
block4_conv2	(Conv2D)	(None, None, None, 512)	2359808
block4_conv3	(Conv2D)	(None, None, None, 512)	2359808
block4_pool	(MaxPooling2D)	(None, None, None, 512)	0
block5_conv1	(Conv2D)	(None, None, None, 512)	2359808
block5_conv2	(Conv2D)	(None, None, None, 512)	2359808
block5_conv3	(Conv2D)	(None, None, None, 512)	2359808
block5_pool	(MaxPooling2D)	(None, None, None, 512)	0

=====

Total params: 14,714,688
Trainable params: 0
Non-trainable params: 14,714,688

Freezing all layers until the fourth from the last

In [34]:

```
conv_base.trainable = True
for layer in conv_base.layers[:-4]:
    layer.trainable = False
```

Fine-tuning the model

In [35]:

```
model.compile(loss="binary_crossentropy",
              optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
```



```
callbacks=callbacks)

Epoch 1/30
188/188 [=====] - 10s 45ms/step - loss: 0.6928 - accuracy: 0.9780 - val_loss: 0.7297 - val_accuracy: 0.9830
Epoch 2/30
188/188 [=====] - 8s 44ms/step - loss: 0.5555 - accuracy: 0.9808 - val_loss: 0.9200 - val_accuracy: 0.9800
Epoch 3/30
188/188 [=====] - 8s 44ms/step - loss: 0.3402 - accuracy: 0.9852 - val_loss: 0.7360 - val_accuracy: 0.9800
Epoch 4/30
188/188 [=====] - 9s 45ms/step - loss: 0.3804 - accuracy: 0.9852 - val_loss: 0.7474 - val_accuracy: 0.9830
Epoch 5/30
188/188 [=====] - 8s 43ms/step - loss: 0.3299 - accuracy: 0.9845 - val_loss: 0.7136 - val_accuracy: 0.9850
Epoch 6/30
188/188 [=====] - 8s 43ms/step - loss: 0.3167 - accuracy: 0.9865 - val_loss: 0.6155 - val_accuracy: 0.9840
Epoch 7/30
188/188 [=====] - 8s 43ms/step - loss: 0.2677 - accuracy: 0.9855 - val_loss: 0.5729 - val_accuracy: 0.9850
Epoch 8/30
188/188 [=====] - 8s 43ms/step - loss: 0.1810 - accuracy: 0.9900 - val_loss: 0.7391 - val_accuracy: 0.9780
Epoch 9/30
188/188 [=====] - 8s 43ms/step - loss: 0.1700 - accuracy: 0.9895 - val_loss: 0.5627 - val_accuracy: 0.9800
Epoch 10/30
188/188 [=====] - 8s 44ms/step - loss: 0.1326 - accuracy: 0.9900 - val_loss: 0.5365 - val_accuracy: 0.9830
Epoch 11/30
188/188 [=====] - 8s 43ms/step - loss: 0.1140 - accuracy: 0.9918 - val_loss: 0.6783 - val_accuracy: 0.9780
Epoch 12/30
188/188 [=====] - 8s 44ms/step - loss: 0.1389 - accuracy: 0.9902 - val_loss: 0.4407 - val_accuracy: 0.9860
Epoch 13/30
188/188 [=====] - 8s 43ms/step - loss: 0.1394 - accuracy: 0.9915 - val_loss: 0.4445 - val_accuracy: 0.9860
Epoch 14/30
188/188 [=====] - 8s 43ms/step - loss: 0.1565 - accuracy: 0.9918 - val_loss: 0.4780 - val_accuracy: 0.9840
Epoch 15/30
188/188 [=====] - 8s 43ms/step - loss: 0.0792 - accuracy: 0.9940 - val_loss: 0.5984 - val_accuracy: 0.9810
Epoch 16/30
188/188 [=====] - 8s 43ms/step - loss: 0.1312 - accuracy: 0.9907 - val_loss: 0.7834 - val_accuracy: 0.9820
Epoch 17/30
188/188 [=====] - 8s 44ms/step - loss: 0.0776 - accuracy: 0.9950 - val_loss: 0.6558 - val_accuracy: 0.9840
```

```

Epoch 18/30
188/188 [=====] - 8s 43ms/step - loss: 0.0899 - accu
racy: 0.9920 - val_loss: 0.6427 - val_accuracy: 0.9840
Epoch 19/30
188/188 [=====] - 8s 43ms/step - loss: 0.0774 - accu
racy: 0.9933 - val_loss: 0.5150 - val_accuracy: 0.9870
Epoch 20/30
188/188 [=====] - 8s 43ms/step - loss: 0.0772 - accu
racy: 0.9940 - val_loss: 0.5283 - val_accuracy: 0.9860
Epoch 21/30
188/188 [=====] - 8s 44ms/step - loss: 0.0716 - accu
racy: 0.9943 - val_loss: 0.5145 - val_accuracy: 0.9830
Epoch 22/30
188/188 [=====] - 8s 43ms/step - loss: 0.0568 - accu
racy: 0.9953 - val_loss: 0.5406 - val_accuracy: 0.9880
Epoch 23/30
188/188 [=====] - 8s 43ms/step - loss: 0.0569 - accu
racy: 0.9957 - val_loss: 0.5038 - val_accuracy: 0.9860
Epoch 24/30
188/188 [=====] - 8s 43ms/step - loss: 0.0556 - accu
racy: 0.9948 - val_loss: 0.6287 - val_accuracy: 0.9860
Epoch 25/30
188/188 [=====] - 8s 43ms/step - loss: 0.0645 - accu
racy: 0.9947 - val_loss: 0.4839 - val_accuracy: 0.9880
Epoch 26/30
188/188 [=====] - 8s 44ms/step - loss: 0.0593 - accu
racy: 0.9953 - val_loss: 0.5203 - val_accuracy: 0.9890
Epoch 27/30
188/188 [=====] - 9s 45ms/step - loss: 0.0666 - accu
racy: 0.9952 - val_loss: 0.4642 - val_accuracy: 0.9850
Epoch 28/30
188/188 [=====] - 8s 44ms/step - loss: 0.0503 - accu
racy: 0.9950 - val_loss: 0.4715 - val_accuracy: 0.9860
Epoch 29/30
188/188 [=====] - 8s 44ms/step - loss: 0.0286 - accu
racy: 0.9975 - val_loss: 0.5109 - val_accuracy: 0.9890
Epoch 30/30
188/188 [=====] - 8s 44ms/step - loss: 0.0466 - accu
racy: 0.9965 - val_loss: 0.6845 - val_accuracy: 0.9810

```

In [36]:

```

model = keras.models.load_model("fine_tuning.keras")
test_loss, test_acc = model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 33ms/step - loss: 0.9459 - accura
cy: 0.9790
Test accuracy: 0.979

```

Summary

Using a pretrained network with training sample of 5000, a validation sample of 500, and a test sample of 500

Downloading the data

```
#!/unzip -qq '/fs/ess/PGS0333/BA_64061_KSU/data/dogs-vs-cats.zip'
```

In [1]:

```
#!/unzip -qq train.zip
```

In [2]:

Copying images to training, validation, and test directories

```
import os, shutil, pathlib
```

In [3]:

```
original_dir = pathlib.Path("train")
new_base_dir = pathlib.Path("cats_vs_dogs_small")

def make_subset(subset_name, start_index, end_index):
    for category in ("cat", "dog"):
        dir = new_base_dir / subset_name / category
        os.makedirs(dir)
        fnames = [f"{category}.{i}.jpg" for i in range(start_index,
end_index)]
        for fname in fnames:
            shutil.copyfile(src=original_dir / fname,
                            dst=dir / fname)

make_subset("train", start_index=0, end_index=4999)
make_subset("validation", start_index=5000, end_index=5499)
make_subset("test", start_index=5500, end_index=5999)
```

Building the model

Instantiating a small convnet for dogs vs. cats classification

```
"""
from tensorflow import keras
from tensorflow.keras import layers
inputs = keras.Input(shape=(180, 180, 3))
x = layers.Rescaling(1./255)(inputs)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
```

In [4]:

```

x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
"""

```

Out[4]:

```

'\nfrom tensorflow import keras\nfrom tensorflow.keras import layers\ninputs
= keras.Input(shape=(180, 180, 3))\nx = layers.Rescaling(1./255)(inputs)\nx =
layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)\nx = layers.M
axPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filters=64, kernel_size=3, act
ivation="relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2
D(filters=128, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling2D
(pool_size=2)(x)\nx = layers.Conv2D(filters=256, kernel_size=3, activation="r
elu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filters=
256, kernel_size=3, activation="relu")(x)\nx = layers.Flatten()(x)\noutputs =
layers.Dense(1, activation="sigmoid")(x)\nmodel = keras.Model(inputs=inputs,
outputs=outputs)\n'

```

In [5]:

```

# model.summary()

```

Configuring the model for training

In [6]:

```

"""
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
"""

```

Out[6]:

```

'\nmodel.compile(loss="binary_crossentropy",\n
optimizer="rmspro
p",\n
metrics=["accuracy"])\n'

```

Data preprocessing

Using image_dataset_from_directory to read images

In [7]:

```

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.utils import image_dataset_from_directory

train_dataset = image_dataset_from_directory(
    new_base_dir / "train",
    image_size=(180, 180),
    batch_size=32)
validation_dataset = image_dataset_from_directory(
    new_base_dir / "validation",
    image_size=(180, 180),
    batch_size=32)
test_dataset = image_dataset_from_directory(
    new_base_dir / "test",

```

```
image_size=(180, 180),
batch_size=32)
```

```
Found 9998 files belonging to 2 classes.
Found 998 files belonging to 2 classes.
Found 998 files belonging to 2 classes.
```

In [8]:

```
"""
import numpy as np
import tensorflow as tf
random_numbers = np.random.normal(size=(1000, 16))
dataset = tf.data.Dataset.from_tensor_slices(random_numbers)
"""
```

Out[8]:

```
'\nimport numpy as np\nimport tensorflow as tf\nrandom_numbers = np.random.normal(size=(1000, 16))\ndataset = tf.data.Dataset.from_tensor_slices(random_numbers)\n'
```

In [9]:

```
"""
for i, element in enumerate(dataset):
    print(element.shape)
    if i >= 2:
        break
"""
```

Out[9]:

```
'\nfor i, element in enumerate(dataset):\n    print(element.shape)\n    if i >= 2:\n        break\n'
```

In [10]:

```
"""
batched_dataset = dataset.batch(32)
for i, element in enumerate(batched_dataset):
    print(element.shape)
    if i >= 2:
        break
"""
```

Out[10]:

```
'\nbatched_dataset = dataset.batch(32)\nfor i, element in enumerate(batched_dataset):\n    print(element.shape)\n    if i >= 2:\n        break\n'
```

In [11]:

```
"""
reshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))
for i, element in enumerate(reshaped_dataset):
    print(element.shape)
    if i >= 2:
        break
"""
```

Out[11]:

```
'\nreshaped_dataset = dataset.map(lambda x: tf.reshape(x, (4, 4)))\nfor i, element in enumerate(reshaped_dataset):\n    print(element.shape)\n    if i >= 2:\n        break\n'
```

Displaying the shapes of the data and labels yielded by the Dataset

In [12]:

```
"""
for data_batch, labels_batch in train_dataset:
    print("data batch shape:", data_batch.shape)
    print("labels batch shape:", labels_batch.shape)
    break
"""
```

Out[12]:

```
'\nfor data_batch, labels_batch in train_dataset:\n    print("data batch shape:", data_batch.shape)\n    print("labels batch shape:", labels_batch.shape)\n    break\n'
```

Fitting the model using a Dataset

In [13]:

```
"""
callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="convnet_from_scratch.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
    callbacks=callbacks)
"""
```

Out[13]:

```
'\ncallbacks = [\n    keras.callbacks.ModelCheckpoint(\n        filepath="convnet_from_scratch.keras",\n        save_best_only=True,\n        monitor="val_loss")\n]\nhistory = model.fit(\n    train_dataset,\n    epochs=30,\n    validation_data=validation_dataset,\n    callbacks=callbacks)\n'
```

Displaying curves of loss and accuracy during training

In [14]:

```
"""
import matplotlib.pyplot as plt
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
"""
```

```
plt.show()
"""
```

Out[14]:

```
'\nimport matplotlib.pyplot as plt\naccuracy = history.history["accuracy"]\nval_accuracy = history.history["val_accuracy"]\nloss = history.history["loss"]\nval_loss = history.history["val_loss"]\nepochs = range(1, len(accuracy) + 1)\nplt.plot(epochs, accuracy, "bo", label="Training accuracy")\nplt.plot(epochs, val_accuracy, "b", label="Validation accuracy")\nplt.title("Training and validation accuracy")\nplt.legend()\nplt.figure()\nplt.plot(epochs, loss, "bo", label="Training loss")\nplt.plot(epochs, val_loss, "b", label="Validation loss")\nplt.title("Training and validation loss")\nplt.legend()\nplt.show()\n'
```

Evaluating the model on the test set

In [15]:

```
"""
test_model = keras.models.load_model("convnet_from_scratch.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
"""
```

Out[15]:

```
'\ntest_model = keras.models.load_model("convnet_from_scratch.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n'
```

Using data augmentation

Define a data augmentation stage to add to an image model

In [16]:

```
"""
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
    ]
)
"""
```

Out[16]:

```
'\ndata_augmentation = keras.Sequential(\n    [\n        layers.RandomFlip("horizontal"),\n        layers.RandomRotation(0.1),\n        layers.RandomZoom(0.2),\n    ]\n)\n'
```

Displaying some randomly augmented training images

In [17]:

```
"""
plt.figure(figsize=(10, 10))
for images, _ in train_dataset.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
```

```

plt.imshow(augmented_images[0].numpy().astype("uint8"))
plt.axis("off")
"""

```

Out[17]:

```

'\nplt.figure(figsize=(10, 10))\nfor images, _ in train_dataset.take(1):\n
for i in range(9):\n    augmented_images = data_augmentation(images)\n
    ax = plt.subplot(3, 3, i + 1)\n    plt.imshow(augmented_images[0].n
umpy().astype("uint8"))\n    plt.axis("off")\n'

```

Defining a new convnet that includes image augmentation and dropout

In [18]:

```

"""
inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = layers.Rescaling(1./255)(x)
x = layers.Conv2D(filters=32, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=64, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=128, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.MaxPooling2D(pool_size=2)(x)
x = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)
x = layers.Flatten()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs=inputs, outputs=outputs)

model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])
"""

```

Out[18]:

```

'\ninputs = keras.Input(shape=(180, 180, 3))\nx = data_augmentation(inputs)\n
x = layers.Rescaling(1./255)(x)\nx = layers.Conv2D(filters=32, kernel_size=3,
activation="relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.C
onv2D(filters=64, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling
2D(pool_size=2)(x)\nx = layers.Conv2D(filters=128, kernel_size=3, activation=
"relu")(x)\nx = layers.MaxPooling2D(pool_size=2)(x)\nx = layers.Conv2D(filter
s=256, kernel_size=3, activation="relu")(x)\nx = layers.MaxPooling2D(pool_siz
e=2)(x)\nx = layers.Conv2D(filters=256, kernel_size=3, activation="relu")(x)\
nx = layers.Flatten()(x)\nx = layers.Dropout(0.5)(x)\noutputs = layers.Dense
(1, activation="sigmoid")(x)\nmodel = keras.Model(inputs=inputs, outputs=outp
uts)\n\nmodel.compile(loss="binary_crossentropy",\n                    optimizer="r
msprop",\n                    metrics=["accuracy"])\n'

```

Training the regularized convnet

In [19]:

```

"""
callbacks = [
    keras.callbacks.ModelCheckpoint(

```



```

        filepath="convnet_from_scratch_with_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=100,
    validation_data=validation_dataset,
    callbacks=callbacks)
"""

```

Out[19]:

```

'\ncallbacks = [\n    keras.callbacks.ModelCheckpoint(\n        filepath="convnet_from_scratch_with_augmentation.keras",\n        save_best_only=True,\n        monitor="val_loss")\n]\nhistory = model.fit(\n    train_dataset,\n    epochs=100,\n    validation_data=validation_dataset,\n    callbacks=callbacks)\n\n'

```

Evaluating the model on the test set

In [20]:

```

"""
test_model = keras.models.load_model(
    "convnet_from_scratch_with_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")
"""

```

Out[20]:

```

'\ntest_model = keras.models.load_model(\n    "convnet_from_scratch_with_augmentation.keras")\ntest_loss, test_acc = test_model.evaluate(test_dataset)\nprint(f"Test accuracy: {test_acc:.3f}")\n\n'

```

Leveraging a pretrained model

Feature extraction with a pretrained model

Instantiating the VGG16 convolutional base

In [21]:

```

from tensorflow import keras # import keras
from tensorflow.keras import layers
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False,
    input_shape=(180, 180, 3))

```

In [22]:

```

conv_base.summary()
Model: "vgg16"

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 180, 180, 3)]	0

block1_conv1 (Conv2D)	(None, 180, 180, 64)	1792
block1_conv2 (Conv2D)	(None, 180, 180, 64)	36928
block1_pool (MaxPooling2D)	(None, 90, 90, 64)	0
block2_conv1 (Conv2D)	(None, 90, 90, 128)	73856
block2_conv2 (Conv2D)	(None, 90, 90, 128)	147584
block2_pool (MaxPooling2D)	(None, 45, 45, 128)	0
block3_conv1 (Conv2D)	(None, 45, 45, 256)	295168
block3_conv2 (Conv2D)	(None, 45, 45, 256)	590080
block3_conv3 (Conv2D)	(None, 45, 45, 256)	590080
block3_pool (MaxPooling2D)	(None, 22, 22, 256)	0
block4_conv1 (Conv2D)	(None, 22, 22, 512)	1180160
block4_conv2 (Conv2D)	(None, 22, 22, 512)	2359808
block4_conv3 (Conv2D)	(None, 22, 22, 512)	2359808
block4_pool (MaxPooling2D)	(None, 11, 11, 512)	0
block5_conv1 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv2 (Conv2D)	(None, 11, 11, 512)	2359808
block5_conv3 (Conv2D)	(None, 11, 11, 512)	2359808
block5_pool (MaxPooling2D)	(None, 5, 5, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

Fast feature extraction without data augmentation

Extracting the VGG16 features and corresponding labels

In [23]:

```
import numpy as np

def get_features_and_labels(dataset):
    all_features = []
    all_labels = []
    for images, labels in dataset:
```

```

        preprocessed_images =
keras.applications.vgg16.preprocess_input(images)
        features = conv_base.predict(preprocessed_images)
        all_features.append(features)
        all_labels.append(labels)
    return np.concatenate(all_features), np.concatenate(all_labels)

train_features, train_labels = get_features_and_labels(train_dataset)
val_features, val_labels = get_features_and_labels(validation_dataset)
test_features, test_labels = get_features_and_labels(test_dataset)

```

In [24]:

```
train_features.shape
```

Out[24]:

```
(9998, 5, 5, 512)
```

Defining and training the densely connected classifier

In [25]:

```

inputs = keras.Input(shape=(5, 5, 512))
x = layers.Flatten()(inputs)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_features, train_labels,
    epochs=20,
    validation_data=(val_features, val_labels),
    callbacks=callbacks)

Epoch 1/20
313/313 [=====] - 1s 3ms/step - loss: 7.6790 - accur
acy: 0.9564 - val_loss: 5.2808 - val_accuracy: 0.9719
Epoch 2/20
313/313 [=====] - 1s 2ms/step - loss: 3.6543 - accur
acy: 0.9779 - val_loss: 5.5756 - val_accuracy: 0.9699
Epoch 3/20
313/313 [=====] - 1s 2ms/step - loss: 1.9265 - accur
acy: 0.9849 - val_loss: 5.5961 - val_accuracy: 0.9749
Epoch 4/20
313/313 [=====] - 1s 2ms/step - loss: 1.5110 - accur
acy: 0.9883 - val_loss: 4.9416 - val_accuracy: 0.9749
Epoch 5/20

```

```

313/313 [=====] - 1s 2ms/step - loss: 0.9044 - accur
acy: 0.9915 - val_loss: 5.3424 - val_accuracy: 0.9729
Epoch 6/20
313/313 [=====] - 1s 2ms/step - loss: 0.7547 - accur
acy: 0.9947 - val_loss: 4.7005 - val_accuracy: 0.9699
Epoch 7/20
313/313 [=====] - 1s 2ms/step - loss: 0.4912 - accur
acy: 0.9940 - val_loss: 5.3495 - val_accuracy: 0.9719
Epoch 8/20
313/313 [=====] - 1s 2ms/step - loss: 0.3394 - accur
acy: 0.9956 - val_loss: 5.9827 - val_accuracy: 0.9749
Epoch 9/20
313/313 [=====] - 1s 2ms/step - loss: 0.3428 - accur
acy: 0.9956 - val_loss: 6.8560 - val_accuracy: 0.9719
Epoch 10/20
313/313 [=====] - 1s 2ms/step - loss: 0.2664 - accur
acy: 0.9964 - val_loss: 7.2837 - val_accuracy: 0.9689
Epoch 11/20
313/313 [=====] - 1s 2ms/step - loss: 0.3664 - accur
acy: 0.9964 - val_loss: 5.8309 - val_accuracy: 0.9780
Epoch 12/20
313/313 [=====] - 1s 2ms/step - loss: 0.2396 - accur
acy: 0.9973 - val_loss: 5.6674 - val_accuracy: 0.9780
Epoch 13/20
313/313 [=====] - 1s 2ms/step - loss: 0.2081 - accur
acy: 0.9976 - val_loss: 5.8744 - val_accuracy: 0.9739
Epoch 14/20
313/313 [=====] - 1s 2ms/step - loss: 0.2907 - accur
acy: 0.9966 - val_loss: 5.7072 - val_accuracy: 0.9729
Epoch 15/20
313/313 [=====] - 1s 2ms/step - loss: 0.1788 - accur
acy: 0.9980 - val_loss: 9.6363 - val_accuracy: 0.9569
Epoch 16/20
313/313 [=====] - 1s 2ms/step - loss: 0.1805 - accur
acy: 0.9971 - val_loss: 7.4781 - val_accuracy: 0.9699
Epoch 17/20
313/313 [=====] - 1s 2ms/step - loss: 0.1058 - accur
acy: 0.9982 - val_loss: 6.7146 - val_accuracy: 0.9729
Epoch 18/20
313/313 [=====] - 1s 2ms/step - loss: 0.0558 - accur
acy: 0.9994 - val_loss: 6.8548 - val_accuracy: 0.9699
Epoch 19/20
313/313 [=====] - 1s 2ms/step - loss: 0.0696 - accur
acy: 0.9990 - val_loss: 7.1342 - val_accuracy: 0.9719
Epoch 20/20
313/313 [=====] - 1s 2ms/step - loss: 0.0474 - accur
acy: 0.9989 - val_loss: 7.2712 - val_accuracy: 0.9709

```

Plotting the results

In [26]:

```


import matplotlib.pyplot as plt
acc = history.history["accuracy"]
val_acc = history.history["val_accuracy"]


```

```

loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(acc) + 1)
plt.plot(epochs, acc, "bo", label="Training accuracy")
plt.plot(epochs, val_acc, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()

```





Feature extraction together with data augmentation

Instantiating and freezing the VGG16 convolutional base

In [27]:

```
conv_base = keras.applications.vgg16.VGG16(
    weights="imagenet",
    include_top=False)
conv_base.trainable = False
```

Printing the list of trainable weights before and after freezing

In [28]:

```
conv_base.trainable = True
print("This is the number of trainable weights "
      "before freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights before freezing the conv base: 26

In [29]:

```
conv_base.trainable = False
print("This is the number of trainable weights "
      "after freezing the conv base:", len(conv_base.trainable_weights))
```

This is the number of trainable weights after freezing the conv base: 0

Adding a data augmentation stage and a classifier to the convolutional base

In [30]:

```
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.2),
```

```

    ]
)

inputs = keras.Input(shape=(180, 180, 3))
x = data_augmentation(inputs)
x = keras.applications.vgg16.preprocess_input(x)
x = conv_base(x)
x = layers.Flatten()(x)
x = layers.Dense(256)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(loss="binary_crossentropy",
              optimizer="rmsprop",
              metrics=["accuracy"])

```

In [31]:

```

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="feature_extraction_with_data_augmentation.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=50,
    validation_data=validation_dataset,
    callbacks=callbacks)

```

Epoch 1/50
313/313 [=====] - 12s 35ms/step - loss: 9.6923 - accuracy: 0.9346 - val_loss: 5.2397 - val_accuracy: 0.9679
Epoch 2/50
313/313 [=====] - 11s 34ms/step - loss: 4.9260 - accuracy: 0.9525 - val_loss: 2.3895 - val_accuracy: 0.9749
Epoch 3/50
313/313 [=====] - 11s 34ms/step - loss: 2.4564 - accuracy: 0.9564 - val_loss: 1.5804 - val_accuracy: 0.9729
Epoch 4/50
313/313 [=====] - 11s 35ms/step - loss: 1.1197 - accuracy: 0.9556 - val_loss: 0.9224 - val_accuracy: 0.9689
Epoch 5/50
313/313 [=====] - 11s 35ms/step - loss: 0.6696 - accuracy: 0.9590 - val_loss: 0.8900 - val_accuracy: 0.9639
Epoch 6/50
313/313 [=====] - 11s 35ms/step - loss: 0.6269 - accuracy: 0.9626 - val_loss: 0.7655 - val_accuracy: 0.9699
Epoch 7/50
313/313 [=====] - 11s 35ms/step - loss: 0.6374 - accuracy: 0.9633 - val_loss: 0.7202 - val_accuracy: 0.9719
Epoch 8/50
313/313 [=====] - 11s 35ms/step - loss: 0.6553 - accuracy: 0.9650 - val_loss: 1.2151 - val_accuracy: 0.9629
Epoch 9/50

313/313 [=====] - 11s 35ms/step - loss: 0.6400 - accuracy: 0.9677 - val_loss: 1.3147 - val_accuracy: 0.9689
Epoch 10/50
313/313 [=====] - 11s 35ms/step - loss: 0.7644 - accuracy: 0.9653 - val_loss: 1.2588 - val_accuracy: 0.9709
Epoch 11/50
313/313 [=====] - 11s 35ms/step - loss: 0.7015 - accuracy: 0.9664 - val_loss: 1.0466 - val_accuracy: 0.9739
Epoch 12/50
313/313 [=====] - 11s 35ms/step - loss: 0.7859 - accuracy: 0.9679 - val_loss: 0.8831 - val_accuracy: 0.9770
Epoch 13/50
313/313 [=====] - 11s 35ms/step - loss: 0.7914 - accuracy: 0.9659 - val_loss: 1.3361 - val_accuracy: 0.9689
Epoch 14/50
313/313 [=====] - 11s 35ms/step - loss: 0.7501 - accuracy: 0.9682 - val_loss: 1.0085 - val_accuracy: 0.9699
Epoch 15/50
313/313 [=====] - 11s 35ms/step - loss: 0.7335 - accuracy: 0.9687 - val_loss: 1.1531 - val_accuracy: 0.9729
Epoch 16/50
313/313 [=====] - 11s 35ms/step - loss: 0.7070 - accuracy: 0.9715 - val_loss: 1.2293 - val_accuracy: 0.9689
Epoch 17/50
313/313 [=====] - 11s 35ms/step - loss: 0.8132 - accuracy: 0.9698 - val_loss: 1.2263 - val_accuracy: 0.9719
Epoch 18/50
313/313 [=====] - 11s 35ms/step - loss: 0.8423 - accuracy: 0.9700 - val_loss: 2.0989 - val_accuracy: 0.9589
Epoch 19/50
313/313 [=====] - 11s 35ms/step - loss: 0.7331 - accuracy: 0.9714 - val_loss: 2.0393 - val_accuracy: 0.9609
Epoch 20/50
313/313 [=====] - 11s 35ms/step - loss: 0.7890 - accuracy: 0.9707 - val_loss: 1.2456 - val_accuracy: 0.9760
Epoch 21/50
313/313 [=====] - 11s 35ms/step - loss: 0.7665 - accuracy: 0.9715 - val_loss: 1.3908 - val_accuracy: 0.9719
Epoch 22/50
313/313 [=====] - 11s 35ms/step - loss: 0.7519 - accuracy: 0.9734 - val_loss: 1.5127 - val_accuracy: 0.9719
Epoch 23/50
313/313 [=====] - 11s 35ms/step - loss: 0.7151 - accuracy: 0.9741 - val_loss: 1.7963 - val_accuracy: 0.9639
Epoch 24/50
313/313 [=====] - 11s 35ms/step - loss: 0.8171 - accuracy: 0.9725 - val_loss: 1.5686 - val_accuracy: 0.9739
Epoch 25/50
313/313 [=====] - 11s 35ms/step - loss: 0.8799 - accuracy: 0.9714 - val_loss: 1.5058 - val_accuracy: 0.9719
Epoch 26/50
313/313 [=====] - 11s 35ms/step - loss: 0.8031 - accuracy: 0.9740 - val_loss: 1.7395 - val_accuracy: 0.9679

Epoch 27/50
313/313 [=====] - 11s 35ms/step - loss: 0.8709 - accuracy: 0.9710 - val_loss: 1.6984 - val_accuracy: 0.9689
Epoch 28/50
313/313 [=====] - 11s 35ms/step - loss: 0.7026 - accuracy: 0.9765 - val_loss: 2.1059 - val_accuracy: 0.9739
Epoch 29/50
313/313 [=====] - 11s 35ms/step - loss: 0.9013 - accuracy: 0.9723 - val_loss: 3.1759 - val_accuracy: 0.9549
Epoch 30/50
313/313 [=====] - 11s 35ms/step - loss: 0.8879 - accuracy: 0.9715 - val_loss: 2.2850 - val_accuracy: 0.9669
Epoch 31/50
313/313 [=====] - 11s 35ms/step - loss: 0.7657 - accuracy: 0.9762 - val_loss: 2.2964 - val_accuracy: 0.9619
Epoch 32/50
313/313 [=====] - 11s 35ms/step - loss: 0.8014 - accuracy: 0.9751 - val_loss: 1.7148 - val_accuracy: 0.9699
Epoch 33/50
313/313 [=====] - 11s 35ms/step - loss: 0.8062 - accuracy: 0.9731 - val_loss: 1.8675 - val_accuracy: 0.9739
Epoch 34/50
313/313 [=====] - 11s 35ms/step - loss: 0.8036 - accuracy: 0.9761 - val_loss: 2.8071 - val_accuracy: 0.9609
Epoch 35/50
313/313 [=====] - 11s 35ms/step - loss: 0.8583 - accuracy: 0.9745 - val_loss: 1.7494 - val_accuracy: 0.9709
Epoch 36/50
313/313 [=====] - 11s 35ms/step - loss: 0.8375 - accuracy: 0.9755 - val_loss: 2.0629 - val_accuracy: 0.9729
Epoch 37/50
313/313 [=====] - 11s 35ms/step - loss: 0.7910 - accuracy: 0.9764 - val_loss: 2.2334 - val_accuracy: 0.9679
Epoch 38/50
313/313 [=====] - 11s 35ms/step - loss: 0.8255 - accuracy: 0.9767 - val_loss: 1.9966 - val_accuracy: 0.9689
Epoch 39/50
313/313 [=====] - 11s 35ms/step - loss: 0.8527 - accuracy: 0.9763 - val_loss: 2.7698 - val_accuracy: 0.9609
Epoch 40/50
313/313 [=====] - 11s 35ms/step - loss: 0.7939 - accuracy: 0.9778 - val_loss: 2.7091 - val_accuracy: 0.9679
Epoch 41/50
313/313 [=====] - 11s 35ms/step - loss: 0.8483 - accuracy: 0.9771 - val_loss: 2.7788 - val_accuracy: 0.9679
Epoch 42/50
313/313 [=====] - 11s 35ms/step - loss: 0.9186 - accuracy: 0.9747 - val_loss: 1.8424 - val_accuracy: 0.9719
Epoch 43/50
313/313 [=====] - 11s 35ms/step - loss: 0.8484 - accuracy: 0.9773 - val_loss: 3.0520 - val_accuracy: 0.9589
Epoch 44/50

```

313/313 [=====] - 11s 35ms/step - loss: 0.8122 - acc
uracy: 0.9754 - val_loss: 1.7815 - val_accuracy: 0.9709
Epoch 45/50
313/313 [=====] - 11s 35ms/step - loss: 0.8498 - acc
uracy: 0.9762 - val_loss: 1.9218 - val_accuracy: 0.9699
Epoch 46/50
313/313 [=====] - 11s 35ms/step - loss: 0.8169 - acc
uracy: 0.9775 - val_loss: 2.2869 - val_accuracy: 0.9689
Epoch 47/50
313/313 [=====] - 11s 35ms/step - loss: 0.9346 - acc
uracy: 0.9753 - val_loss: 2.1793 - val_accuracy: 0.9699
Epoch 48/50
313/313 [=====] - 11s 35ms/step - loss: 0.8180 - acc
uracy: 0.9790 - val_loss: 4.1362 - val_accuracy: 0.9539
Epoch 49/50
313/313 [=====] - 11s 35ms/step - loss: 0.8713 - acc
uracy: 0.9769 - val_loss: 2.7774 - val_accuracy: 0.9619
Epoch 50/50
313/313 [=====] - 11s 35ms/step - loss: 0.9165 - acc
uracy: 0.9767 - val_loss: 2.1626 - val_accuracy: 0.9739

```

Evaluating the model on the test set

In [32]:

```

test_model = keras.models.load_model(
    "feature_extraction_with_data_augmentation.keras")
test_loss, test_acc = test_model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 30ms/step - loss: 0.4685 - accura
cy: 0.9850
Test accuracy: 0.985

```

Fine-tuning a pretrained model

In [33]:

```

conv_base.summary()
Model: "vgg16"

```

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584

block2_pool	(MaxPooling2D)	(None, None, None, 128)	0
block3_conv1	(Conv2D)	(None, None, None, 256)	295168
block3_conv2	(Conv2D)	(None, None, None, 256)	590080
block3_conv3	(Conv2D)	(None, None, None, 256)	590080
block3_pool	(MaxPooling2D)	(None, None, None, 256)	0
block4_conv1	(Conv2D)	(None, None, None, 512)	1180160
block4_conv2	(Conv2D)	(None, None, None, 512)	2359808
block4_conv3	(Conv2D)	(None, None, None, 512)	2359808
block4_pool	(MaxPooling2D)	(None, None, None, 512)	0
block5_conv1	(Conv2D)	(None, None, None, 512)	2359808
block5_conv2	(Conv2D)	(None, None, None, 512)	2359808
block5_conv3	(Conv2D)	(None, None, None, 512)	2359808
block5_pool	(MaxPooling2D)	(None, None, None, 512)	0

=====

Total params: 14,714,688
Trainable params: 0
Non-trainable params: 14,714,688

Freezing all layers until the fourth from the last

In [34]:

```
conv_base.trainable = True
for layer in conv_base.layers[:-4]:
    layer.trainable = False
```

Fine-tuning the model

In [35]:

```
model.compile(loss="binary_crossentropy",
              optimizer=keras.optimizers.RMSprop(learning_rate=1e-5),
              metrics=["accuracy"])

callbacks = [
    keras.callbacks.ModelCheckpoint(
        filepath="fine_tuning.keras",
        save_best_only=True,
        monitor="val_loss")
]
history = model.fit(
    train_dataset,
    epochs=30,
    validation_data=validation_dataset,
```

```
callbacks=callbacks)

Epoch 1/30
313/313 [=====] - 14s 39ms/step - loss: 0.7166 - acc
uracy: 0.9767 - val_loss: 1.6411 - val_accuracy: 0.9749
Epoch 2/30
313/313 [=====] - 12s 38ms/step - loss: 0.4809 - acc
uracy: 0.9794 - val_loss: 1.8219 - val_accuracy: 0.9729
Epoch 3/30
313/313 [=====] - 12s 38ms/step - loss: 0.4577 - acc
uracy: 0.9802 - val_loss: 1.4528 - val_accuracy: 0.9709
Epoch 4/30
313/313 [=====] - 12s 38ms/step - loss: 0.2875 - acc
uracy: 0.9846 - val_loss: 1.2672 - val_accuracy: 0.9719
Epoch 5/30
313/313 [=====] - 12s 38ms/step - loss: 0.2989 - acc
uracy: 0.9837 - val_loss: 0.9301 - val_accuracy: 0.9780
Epoch 6/30
313/313 [=====] - 12s 38ms/step - loss: 0.2356 - acc
uracy: 0.9844 - val_loss: 0.9989 - val_accuracy: 0.9719
Epoch 7/30
313/313 [=====] - 12s 39ms/step - loss: 0.2032 - acc
uracy: 0.9842 - val_loss: 1.1434 - val_accuracy: 0.9739
Epoch 8/30
313/313 [=====] - 12s 38ms/step - loss: 0.1712 - acc
uracy: 0.9864 - val_loss: 1.0423 - val_accuracy: 0.9729
Epoch 9/30
313/313 [=====] - 12s 39ms/step - loss: 0.1669 - acc
uracy: 0.9870 - val_loss: 0.6818 - val_accuracy: 0.9719
Epoch 10/30
313/313 [=====] - 12s 38ms/step - loss: 0.1539 - acc
uracy: 0.9859 - val_loss: 0.8058 - val_accuracy: 0.9729
Epoch 11/30
313/313 [=====] - 12s 39ms/step - loss: 0.1168 - acc
uracy: 0.9881 - val_loss: 0.7075 - val_accuracy: 0.9790
Epoch 12/30
313/313 [=====] - 12s 38ms/step - loss: 0.1092 - acc
uracy: 0.9900 - val_loss: 0.7504 - val_accuracy: 0.9770
Epoch 13/30
313/313 [=====] - 12s 38ms/step - loss: 0.0905 - acc
uracy: 0.9914 - val_loss: 0.9088 - val_accuracy: 0.9719
Epoch 14/30
313/313 [=====] - 12s 39ms/step - loss: 0.1122 - acc
uracy: 0.9895 - val_loss: 0.5597 - val_accuracy: 0.9760
Epoch 15/30
313/313 [=====] - 12s 38ms/step - loss: 0.0837 - acc
uracy: 0.9897 - val_loss: 0.9406 - val_accuracy: 0.9709
Epoch 16/30
313/313 [=====] - 12s 39ms/step - loss: 0.0726 - acc
uracy: 0.9915 - val_loss: 0.9343 - val_accuracy: 0.9669
Epoch 17/30
313/313 [=====] - 12s 39ms/step - loss: 0.0771 - acc
uracy: 0.9914 - val_loss: 1.0177 - val_accuracy: 0.9689
```

```

Epoch 18/30
313/313 [=====] - 12s 38ms/step - loss: 0.0694 - acc
uracy: 0.9932 - val_loss: 0.7923 - val_accuracy: 0.9749
Epoch 19/30
313/313 [=====] - 12s 39ms/step - loss: 0.0497 - acc
uracy: 0.9949 - val_loss: 0.7644 - val_accuracy: 0.9749
Epoch 20/30
313/313 [=====] - 12s 39ms/step - loss: 0.0601 - acc
uracy: 0.9934 - val_loss: 0.7196 - val_accuracy: 0.9749
Epoch 21/30
313/313 [=====] - 12s 38ms/step - loss: 0.0532 - acc
uracy: 0.9943 - val_loss: 0.6580 - val_accuracy: 0.9760
Epoch 22/30
313/313 [=====] - 12s 39ms/step - loss: 0.0531 - acc
uracy: 0.9933 - val_loss: 0.9021 - val_accuracy: 0.9739
Epoch 23/30
313/313 [=====] - 12s 38ms/step - loss: 0.0663 - acc
uracy: 0.9918 - val_loss: 0.6082 - val_accuracy: 0.9770
Epoch 24/30
313/313 [=====] - 12s 38ms/step - loss: 0.0533 - acc
uracy: 0.9943 - val_loss: 0.6947 - val_accuracy: 0.9770
Epoch 25/30
313/313 [=====] - 12s 38ms/step - loss: 0.0404 - acc
uracy: 0.9949 - val_loss: 0.5410 - val_accuracy: 0.9760
Epoch 26/30
313/313 [=====] - 12s 39ms/step - loss: 0.0479 - acc
uracy: 0.9942 - val_loss: 0.7172 - val_accuracy: 0.9760
Epoch 27/30
313/313 [=====] - 12s 39ms/step - loss: 0.0333 - acc
uracy: 0.9953 - val_loss: 0.5674 - val_accuracy: 0.9820
Epoch 28/30
313/313 [=====] - 12s 38ms/step - loss: 0.0354 - acc
uracy: 0.9957 - val_loss: 0.5682 - val_accuracy: 0.9770
Epoch 29/30
313/313 [=====] - 12s 39ms/step - loss: 0.0309 - acc
uracy: 0.9961 - val_loss: 0.6533 - val_accuracy: 0.9790
Epoch 30/30
313/313 [=====] - 12s 39ms/step - loss: 0.0487 - acc
uracy: 0.9945 - val_loss: 0.6744 - val_accuracy: 0.9770

```

In [36]:

```

model = keras.models.load_model("fine_tuning.keras")
test_loss, test_acc = model.evaluate(test_dataset)
print(f"Test accuracy: {test_acc:.3f}")

32/32 [=====] - 1s 30ms/step - loss: 0.3557 - accura
cy: 0.9810
Test accuracy: 0.981

```

Summary